

CS 1217

Lecture 4 – Process Creation, Termination

Logistics

- Assignment # 2 goes out this week
- Mandatory lab hours, starting this week
 - Tuesday 6:30 – 8 pm
 - Will send out the room details shortly

Recap: pmap

- Information about memory mappings of process

```
cs304@cs304-devel:~$ pmap 2002
2002:  bash
0000560445ea5000    1040K r-x--  bash
00005604461a8000     16K r---  bash
00005604461ac000     36K rw---  bash
00005604461b5000     40K rw---  [ anon ]
00005604480ff000   1604K rw---  [ anon ]

00007fb4b8a15000   1948K r-x--  libc-2.27.so
00007fb4b8bfc000   2048K ----  libc-2.27.so
00007fb4b8dfc000     16K r---  libc-2.27.so
00007fb4b8e00000      8K rw---  libc-2.27.so

00007fb4b945d000      4K rw---  [ anon ]
00007fff083df000   132K rw---  [ stack ]
```

Aside /proc

```
root@cs304-devel:/proc# ls
1      1097  1163  1407  1585  1648  20    269  297  525  694  asound  kallsyms  schedstat
10     11    1194  1415  1586  165   21    27   298  537  7    buddyinfo kcore     scsi
1000   1100  12    1420  1588  1685  2151  270  299  543  766  bus     keys      self
1001   1102  1219  1431  1589  1686  219   271  30   544  78   cgroups  key-users slabinfo
1012   1104  1222  1480  1594  1690  22    272  308  549  780  cmdline kmsg      softirqs
1014   1107  1223  1484  1596  17    2222  2764 309  562  788  consoles kpagecgroup stat
1016   1109  1236  1486  16    1724  2232  2779 31   568  79   cpuinfo  kpagecount swaps
1022   1110  1240  1488  1603  1740  2278  2781 312  569  8    crypto   kpageflags sys
1029   1111  1242  149   1605  1760  23    2782 314  570  80   devices  loadavg   sysrq-trigger
1038   1112  1247  15    1606  1775  2384  2783 318  579  81   diskstats locks      sysvipc
1042   1115  1250  1503  1614  1789  24    2794 32   581  82   dma      mdstat    thread-self
1047   1117  13    1507  1615  1798  2409  28   334  583  83   driver   meminfo   timer_list
1049   1121  1358  1518  1616  18    242   2814 336  584  84   execdomains misc       tty
1052   1122  1359  1548  1618  1839  2495  282  34   585  85   fb        modules   uptime
1054   1123  1363  1554  1622  187   25    285  348  587  89   filesystems mounts     version
1070   1126  1378  1557  1623  188   26    286  35   6    9    fs        mtrr      version_signature
1073   1131  1381  1566  1625  19    263   289  36   621  98   interrupts net        vmallocinfo
1075   1134  1386  1571  1627  1911  264   29   4    624  982  iomem    pagetypeinfo vmstat
1078   1138  1389  1575  163   1984  267   291  411  692  991  ioports  partitions zoneinfo
1093   115   14    1579  164   2     2671  295  5    693  acpi     irq       sched_debug
```

- Can you read from it?
- Can you write to it?
- echo 3 > /proc/sys/vm/drop_caches
- (alright, I lied before)

cpuinfo and meminfo

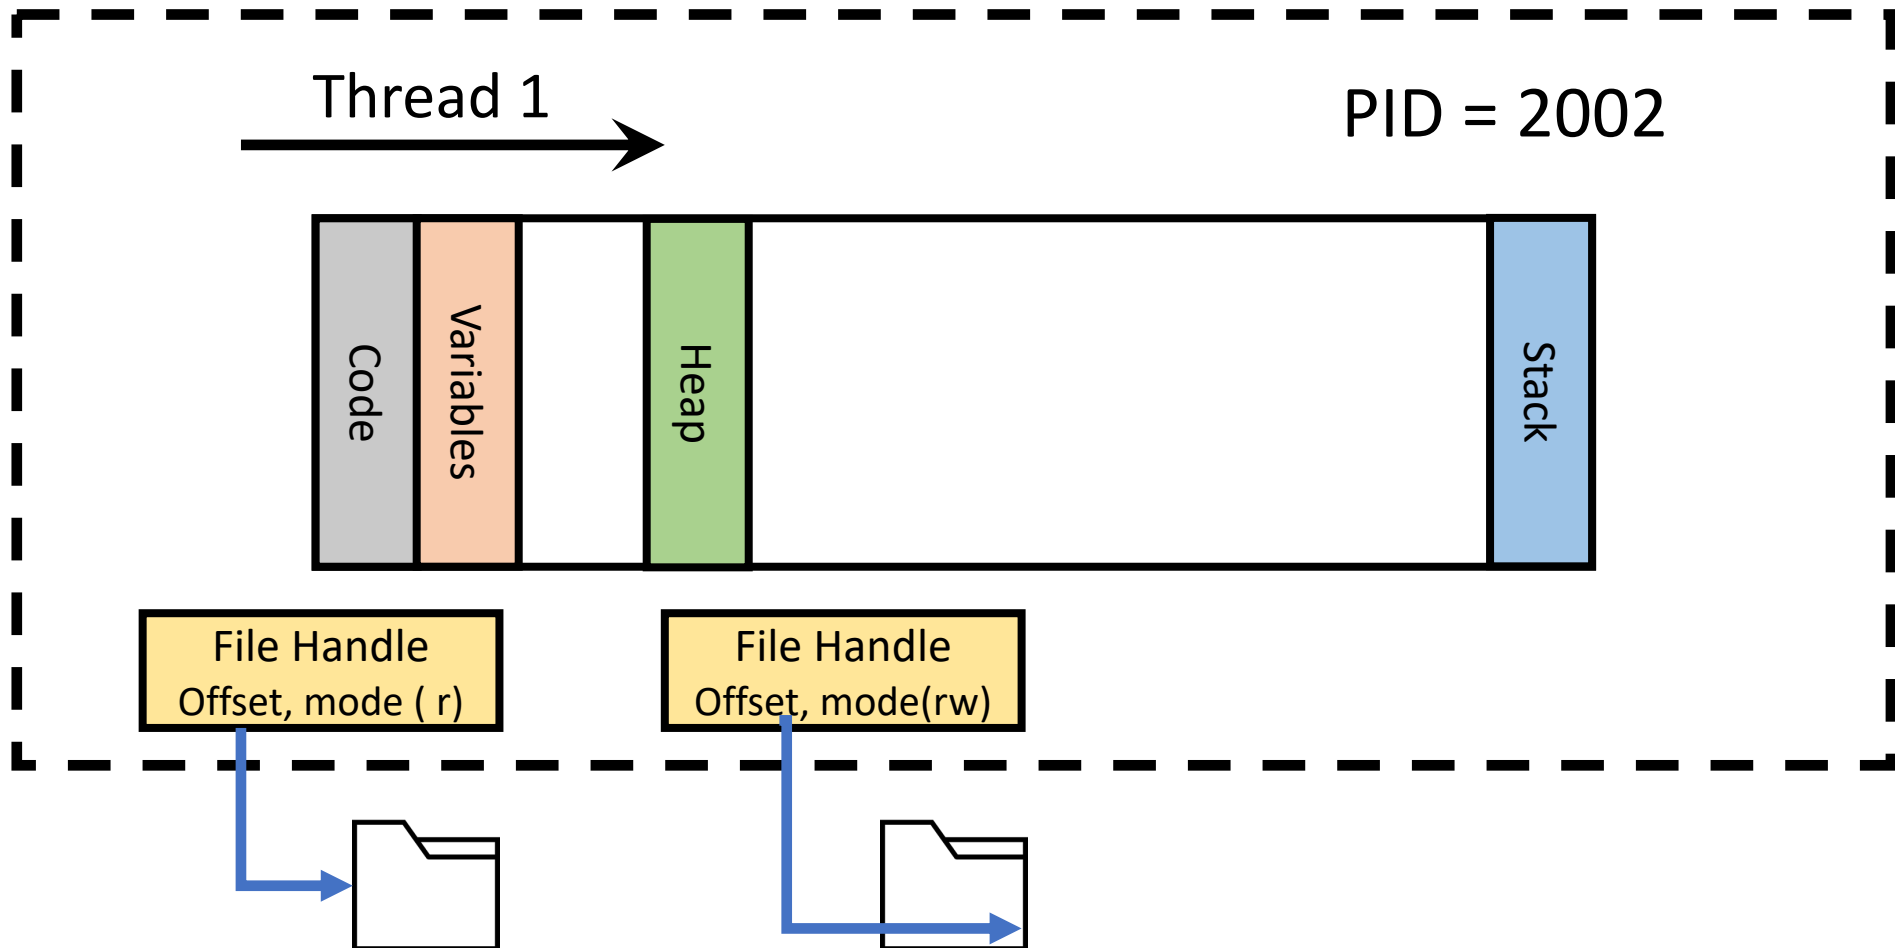
```
cs304@cs304-devel:~$ cat /proc/cpuinfo | more
processor       : 0
vendor_id      : GenuineIntel
cpu family     : 6
model          : 61
model name     : Intel(R) Core(TM) i5-5257U CPU @ 2.70GHz
stepping       : 4
cpu MHz        : 2699.998
cache size     : 3072 KB
physical id    : 0
siblings       : 1
core id        : 0
cpu cores      : 1
apicid         : 0
initial apicid : 0
fpu            : yes
fpu_exception  : yes
cpuid level    : 20
wp             : yes
```

```
cs304@cs304-devel:~$ cat /proc/meminfo
```

```
MemTotal:      2041304 kB
MemFree:       193196 kB
MemAvailable:   925720 kB
Buffers:       54232 kB
Cached:        790040 kB
SwapCached:    0 kB
Active:        1335996 kB
Inactive:      354308 kB
Active(anon):  846872 kB
Inactive(anon): 5368 kB
Active(file):  489124 kB
Inactive(file): 348940 kB
Unevictable:   16 kB
Mlocked:       16 kB
SwapTotal:     728520 kB
SwapFree:      728520 kB
Dirty:         92 kB
Writeback:     0 kB
AnonPages:     846056 kB
Mapped:        225512 kB
Shmem:         6212 kB
Slab:          83600 kB
SReclaimable:  52028 kB
SUnreclaim:    31572 kB
KernelStack:   7820 kB
PageTables:    38140 kB
NFS Unstable:   0 kB
```

The Process Abstraction

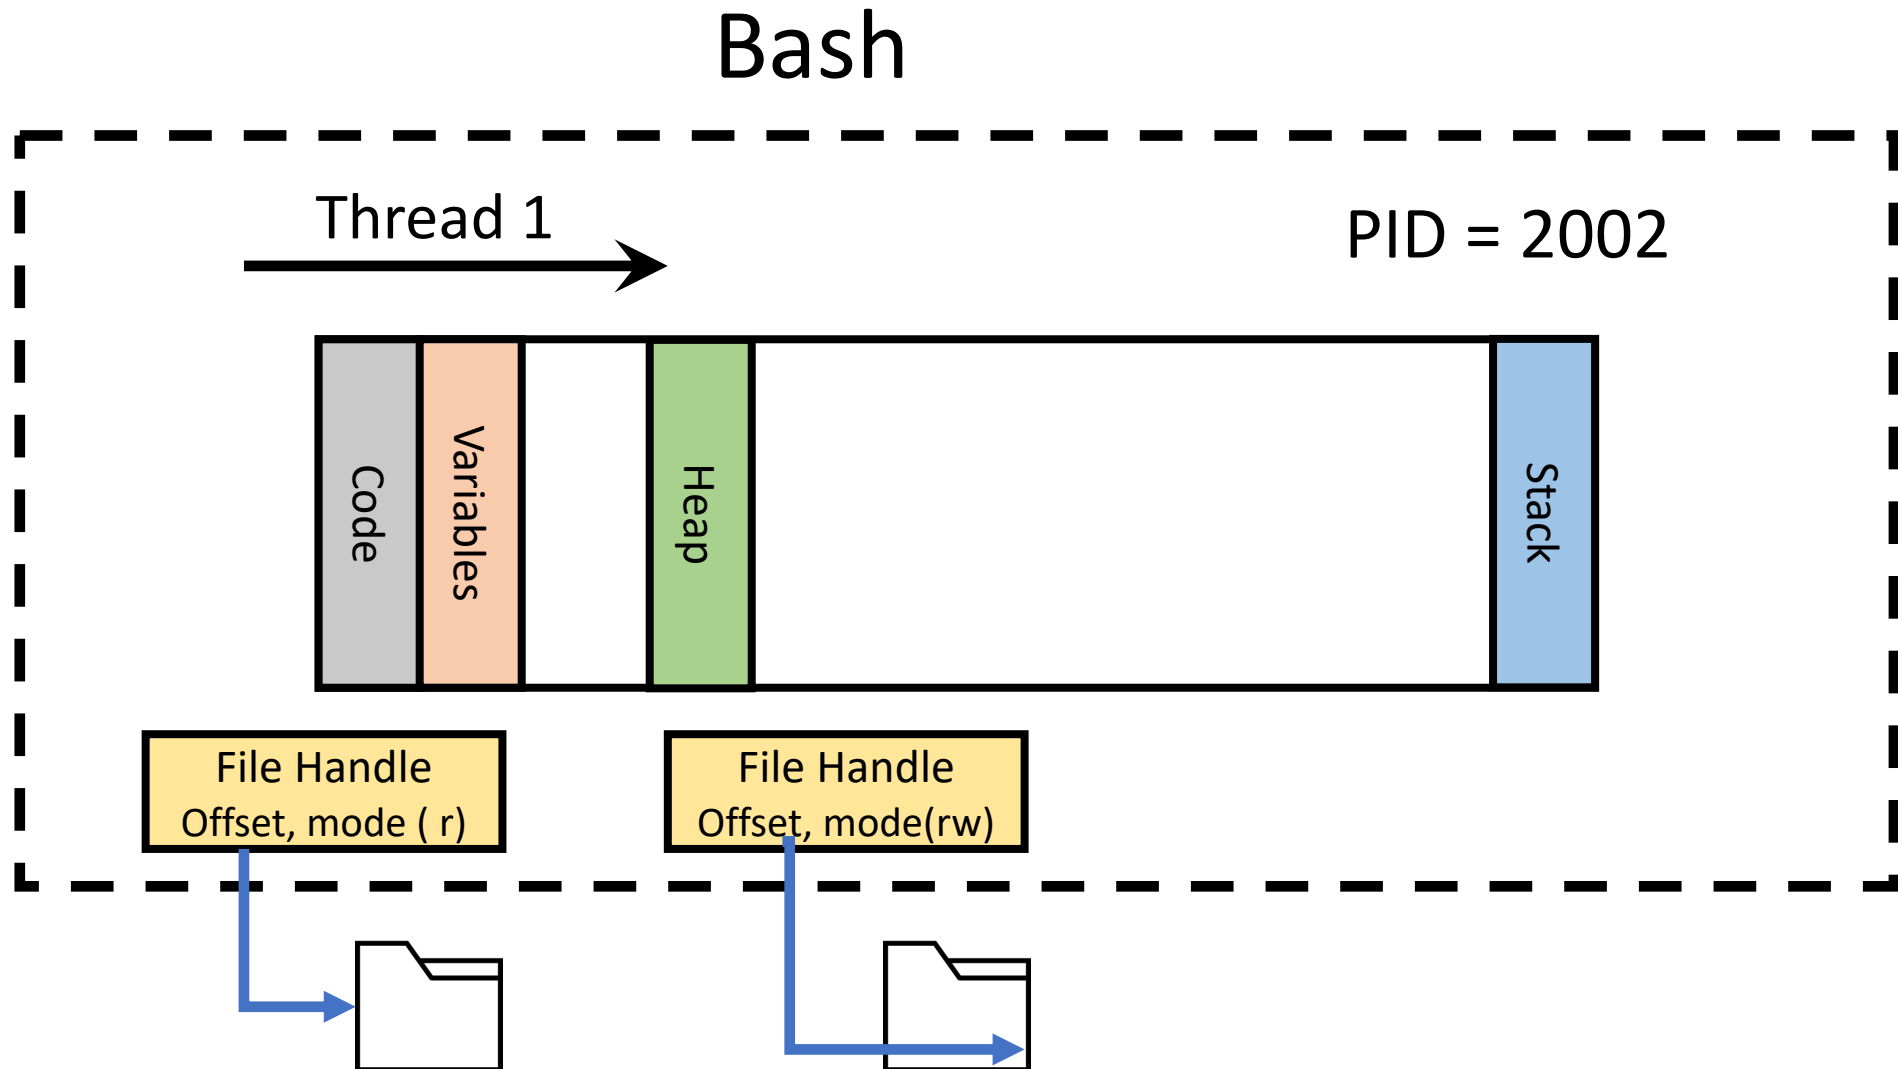
Bash



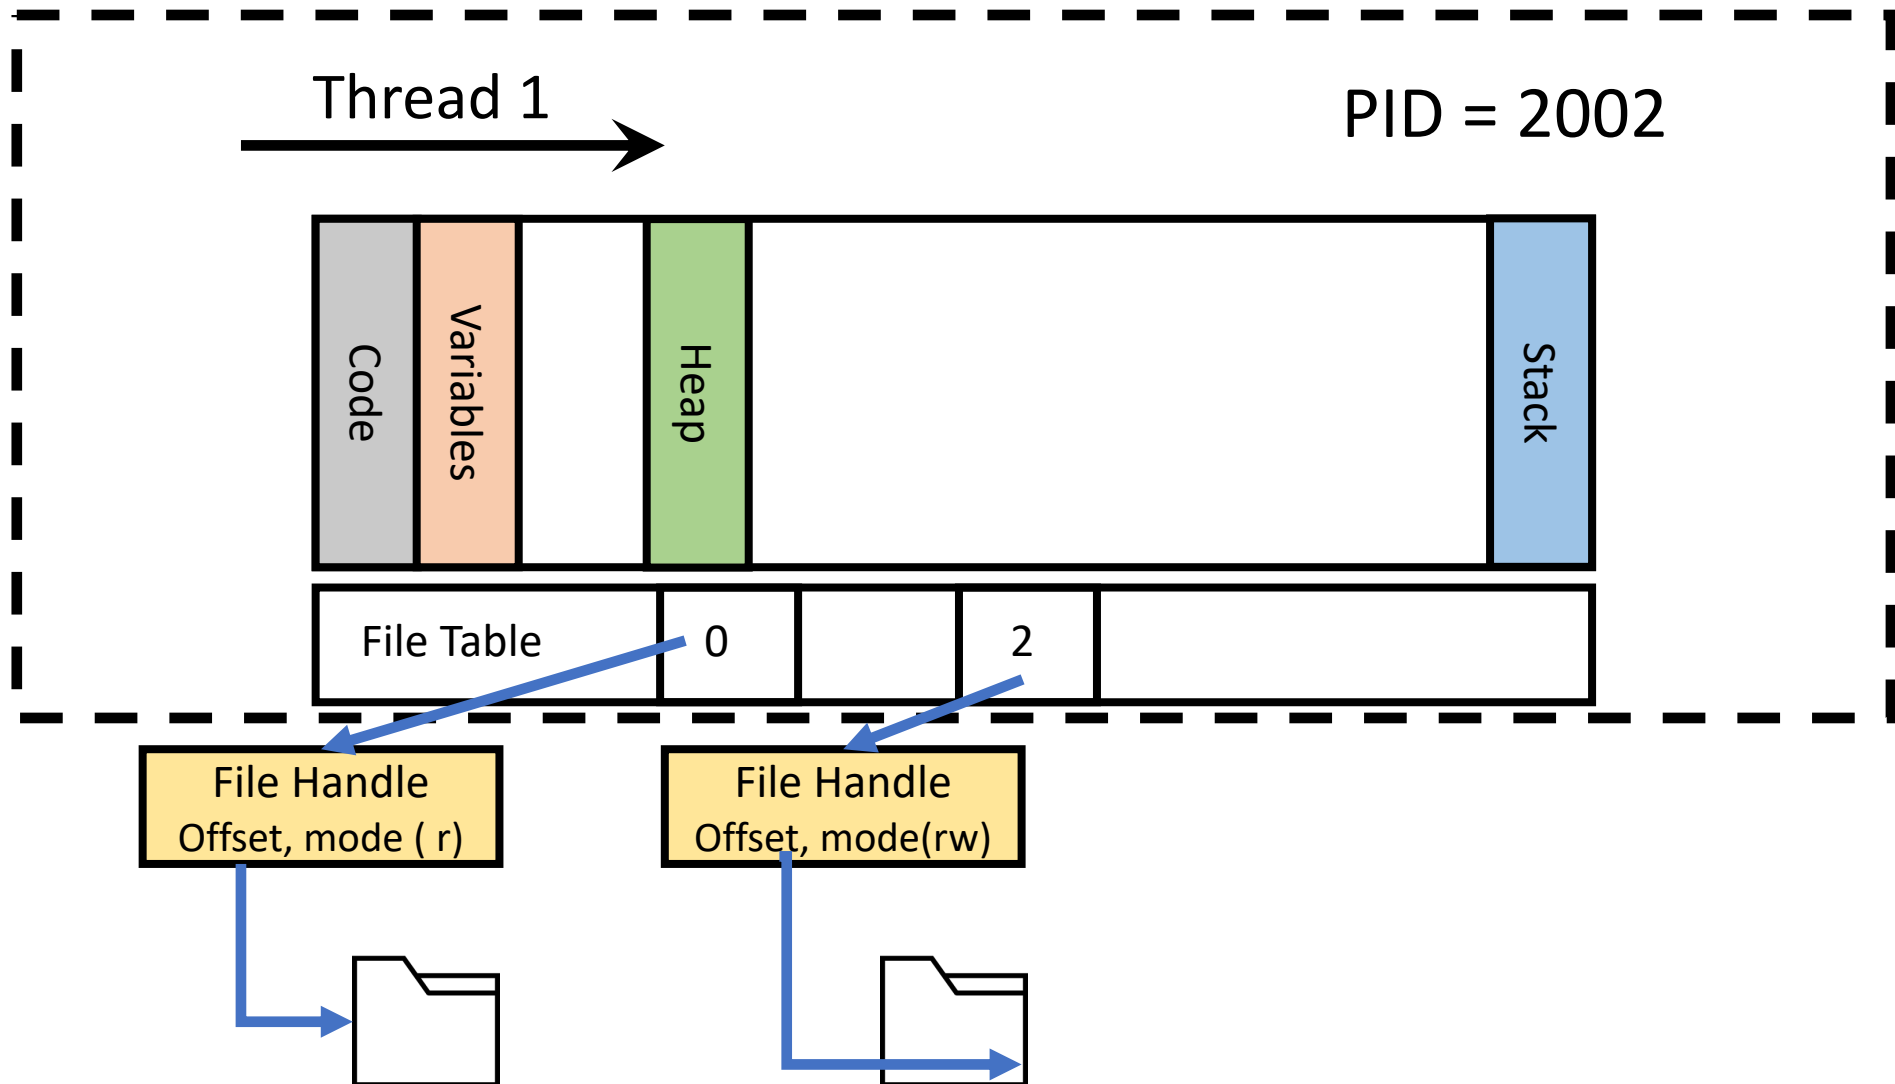
Aside: lsof example

```
cs304@cs304-devel:~$ lsof -p 2002 | more
COMMAND  PID  USER  FD  TYPE  DEVICE  SIZE/OFF  NODE  NAME
bash     2002  cs304  cwd   DIR    8,1      4096  426011 /home/cs304
bash     2002  cs304  rtd   DIR    8,1      4096     2  /
bash     2002  cs304  txt   REG    8,1    1113504  262153 /bin/bash
bash     2002  cs304  mem   REG    8,1     47568  268654 /lib/x86_64-linux-gnu/libnss_files-2.27.so
bash     2002  cs304  mem   REG    8,1     97176  268648 /lib/x86_64-linux-gnu/libnsl-2.27.so
bash     2002  cs304  mem   REG    8,1     47576  268665 /lib/x86_64-linux-gnu/libnss_nis-2.27.so
bash     2002  cs304  mem   REG    8,1     39744  268650 /lib/x86_64-linux-gnu/libnss_compat-2.27.so
bash     2002  cs304  mem   REG    8,1   10281936  794322 /usr/lib/locale/locale-archive
bash     2002  cs304  mem   REG    8,1   2030544  268564 /lib/x86_64-linux-gnu/libc-2.27.so
bash     2002  cs304  mem   REG    8,1     14560  268587 /lib/x86_64-linux-gnu/libdl-2.27.so
bash     2002  cs304  mem   REG    8,1    170784  268722 /lib/x86_64-linux-gnu/libtinfo.so.5.9
bash     2002  cs304  mem   REG    8,1    170960  268536 /lib/x86_64-linux-gnu/ld-2.27.so
bash     2002  cs304  mem   REG    8,1     26376  139948 /usr/lib/x86_64-linux-gnu/gconv/gconv-modules
.cache
bash     2002  cs304   0u   CHR   136,1      0t0     4  /dev/pts/1
bash     2002  cs304   1u   CHR   136,1      0t0     4  /dev/pts/1
bash     2002  cs304   2u   CHR   136,1      0t0     4  /dev/pts/1
bash     2002  cs304  255u  CHR   136,1      0t0     4  /dev/pts/1
```

File Handles



File Tables and File Handles



A Bit about File Tables

- file descriptor → file handle.
- file handle → file object.
- file object → blocks on disk

- What are these "links" called?

- Are these indirections useful? Why?

Why do this?

- The additional level of indirection allows certain pieces of state to be shared separately.
- **File descriptors** are private to each process.
- **File handles** are private to each process but shared after process creation.
 - **File handles** store the current file **offset**, or the position in the file that the next read will come from or write will go to.
 - File handles can be **deliberately** shared between two processes.
- **File objects** hold some file state and may be **shared transparently** between many processes

Where do processes come from?

- The first process is “init”
 - The process that is created at bootup, every other process descends from it
- The OS provides support for existing processes to “create” new processes
 - Parent-child relationships between processes
 - Hence, process trees, check **pstree**
- New process would need the same abstractions as parent
 - CPU state (Threads)
 - Address Spaces (Memory)
 - Open file handles, if need be

The Lifecycle of a Process

- Birth (Create)
 - Who creates a process? What does it mean to create a process?
 - OS Interfaces for process creation
- Death (Destroy)
 - Who destroys a process? What does it mean to destroy a process?
 - What are the OS interfaces to do so
- Wait()
 - Wait for something to happen or complete
- Other Interfaces to a Process
 - Mechanisms to suspend or resume process execution
 - Status query – current state, runtime etc.

Process States

```
27 struct context {
28     uint edi;
29     uint esi;
30     uint ebx;
31     uint ebp;
32     uint eip;
33 };
```


I want to capture *thread* state –
what do these things represent?

```
35 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
36
37 // Per-process state
38 struct proc {
39     uint sz; // Size of process memory (bytes)
40     pde_t* pgdir; // Page table
41     char *kstack; // Bottom of kernel stack for this process
42     enum procstate state; // Process state
43     int pid; // Process ID
44     struct proc *parent; // Parent process
45     struct trapframe *tf; // Trap frame for current syscall
46     struct context *context; // swtch() here to run process
47     void *chan; // If non-zero, sleeping on chan
48     int killed; // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd; // Current directory
51     char name[16]; // Process name (debugging)
52 };
```

Process Birth: fork()

- fork () allows a process to create a new process
- What does that mean?
- fork() copies the caller process.
- fork() copies the address space.
- fork() copies the process file table.

Process Creation : fork()



Process creation : fork()

```
int returnCode = fork();

if (returnCode == 0) {
    // child (new process)
    printf("hello, I am child (pid:%d)\n", (int) getpid());
} else {
    // parent goes down this path (original process)
    printf("hello, I am parent of %d (pid:%d)\n",
        _____, (int) getpid());
}
```

- How many times do you think fork() returns?
- What are the return values?