

CS 1217

Lecture 9 – Context Switch wrap; Process States, Intro to Scheduling

Software exceptions

- Software exception types
 - **Faults:** These can be corrected, and the program may continue as if nothing happened.
 - **Aborts:** Some severe unrecoverable error; leads to code dumps/stack traces
- Divide instruction implementation raises exception if denominator = 0

EXCEPTION PROCESSING

The CPU's response to unusual internal or external conditions.

EXCEPTION VECTORS:

Number	Addr	Use	CLKs	Reads	Writes
Dec	Hex	Hex			
0	00	000000	(Reset SSP; see note below)		
1	01	000004	RESET* 40.6.0		
2	02	000008	Bus Error (BERR*) 50.4.7		
3	03	00000C	Address Error 50.4.7		
4	04	000010	Illegal Instruction 34.4.3		
5	05	000014	Divide by zero 42.5.4		
6	06	000018	CHK operand out of bounds		
7	07	00001C	TRAPV when V=1 34.4.3		
8	08	000020	Privilege violation 34.4.3		
9	09	000024	Trace 34.4.3		
10	0A	000028	Line 1010 emulator 34.4.3		
11	0B	00002C	Line 1111 emulator 34.4.3		
12	0C	000030	(reserved)		
13	0D	000034	(reserved)		
14	0E	000038	(reserved)		
15	0F	00003C	Uninitialized irpt 44.5.3		

Motorola 68000

Batch Processing

- As time progressed: More users; multiple, interactive jobs
- Form a queue, send another job to execute when one finishes



Problems with Batch Computing?

- **Inefficiency!**

- Usage of slower parts of the system will cause the CPU to stall waiting for them to finish
- CPUs are very fast, everything else in the system is comparatively slower
 - How do you increase CPU utilization
- How do you support multiple, concurrent, interactive users?

Supporting Multiple Users




Time scale : seconds

Co-Operative Scheduling

- While a program is doing something slow (e.g. reading the disk), the kernel can find some other process to run, while this process is waiting for the operation to complete.
- What does this depend on?
- What can do wrong here?

Context Switching



Time scale : micro seconds

Why does this work?

- Human perception has some limits
 - 15 ms (or so)

Implementing Context Switches

- Who has to implement the switching between processes/threads?
- How does the OS get to run?
- What if none of the processes do any of those things?

Timer Interrupts

- **Timer interrupts** generated by a timer device ensure that the operating system regains control of the system at regular intervals
- Form the basis of **pre-emptive scheduling**: the OS doesn't have to wait for a process to let it run; can take control and preempt it
- Timer interrupts are handled the same way as other interrupts

Saving Thread State

- A thread of the process is executing at any given time
- Threads can be pre-empted : when they restart, should have the exact state when they were pre-empted
 - It should appear to the thread that nothing has happened, the execution was not interrupted
- What does the thread state consist of?
 - Registers
 - Address space
- Who saves the state?

Trap Frame

- Saving thread state is the **first thing** that happens when the interrupt service routine is triggered (**Why?**)
- Saved state is also known as a **trap frame**
- The kernel switches to a different stack – why?

```
0600 // Layout of the trap frame built on the stack by the
0601 // hardware and by trapasm.S, and passed to trap().
0602 struct trapframe {
0603     // registers as pushed by pusha
0604     uint edi;
0605     uint esi;
0606     uint ebp;
0607     uint oesp;        // useless & ignored
0608     uint ebx;
0609     uint edx;
0610     uint ecx;
0611     uint eax;
0612
0613     // rest of trap frame
0614     ushort gs;
0615     ushort padding1;
0616     ushort fs;
0617     ushort padding2;
0618     ushort es;
0619     ushort padding3;
0620     ushort ds;
0621     ushort padding4;
0622     uint trapno;
0623
0624     // below here defined by x86 hardware
0625     uint err;
0626     uint eip;
0627     ushort cs;
0628     ushort padding5;
0629     uint eflags;
0630
0631     // below here only when crossing rings, such as from user to kernel
0632     uint esp;
0633     ushort ss;
0634     ushort padding6;
0635 };
```

xv6 trap frame

xv6 Context Switch

```
1  # Context switch
2  #
3  # void swtch(struct context **old, struct context *new);
4  #
5  # Save the current registers on the stack, creating
6  # a struct context, and save its address in *old.
7  # Switch stacks to new and pop previously-saved registers.
8
9  .globl swtch
10 swtch:
11     movl 4(%esp), %eax
12     movl 8(%esp), %edx
13
14     # Save old callee-saved registers
15     pushl %ebp
16     pushl %ebx
17     pushl %esi
18     pushl %edi
19
20     # Switch stacks
21     movl %esp, (%eax)
22     movl %edx, %esp
23
24     # Load new callee-saved registers
25     popl %edi
26     popl %esi
27     popl %ebx
28     popl %ebp
29     ret
```


Context Switch Frequency

- Context switches have associated overhead, both in *space* and in *time*
 - Why?
- How does that affect the rate at which context switches can happen?

Process States

- Processes can be thought about as being in different “states”
- **Running** : Currently executing instructions on the CPU
- **Ready** : Not currently executing, but can; is capable of getting restarted
- **Waiting/Blocked/Sleeping**: not executing instructions and not able to be restarted until *some* event occurs

Process States



Scheduling

- **Scheduling** is the process of choosing the next process (or processes) to run on the CPU (or CPUs).
- Why is scheduling required?

When to Schedule Processes?

- When a process **voluntarily** gives up the CPU by calling `yield()`
- When a process makes a blocking **system call** and must sleep until the call completes.
- When a thread **exits**
- When the **kernel decides** that a thread has run for long enough
- Which one is co-operative? Which one is pre-emptive?

Quick note on yield()

- `yield()` can be a useful way of allowing a ***well-behaved*** thread to tell the CPU that it has no more useful work to do.