

CS 1217

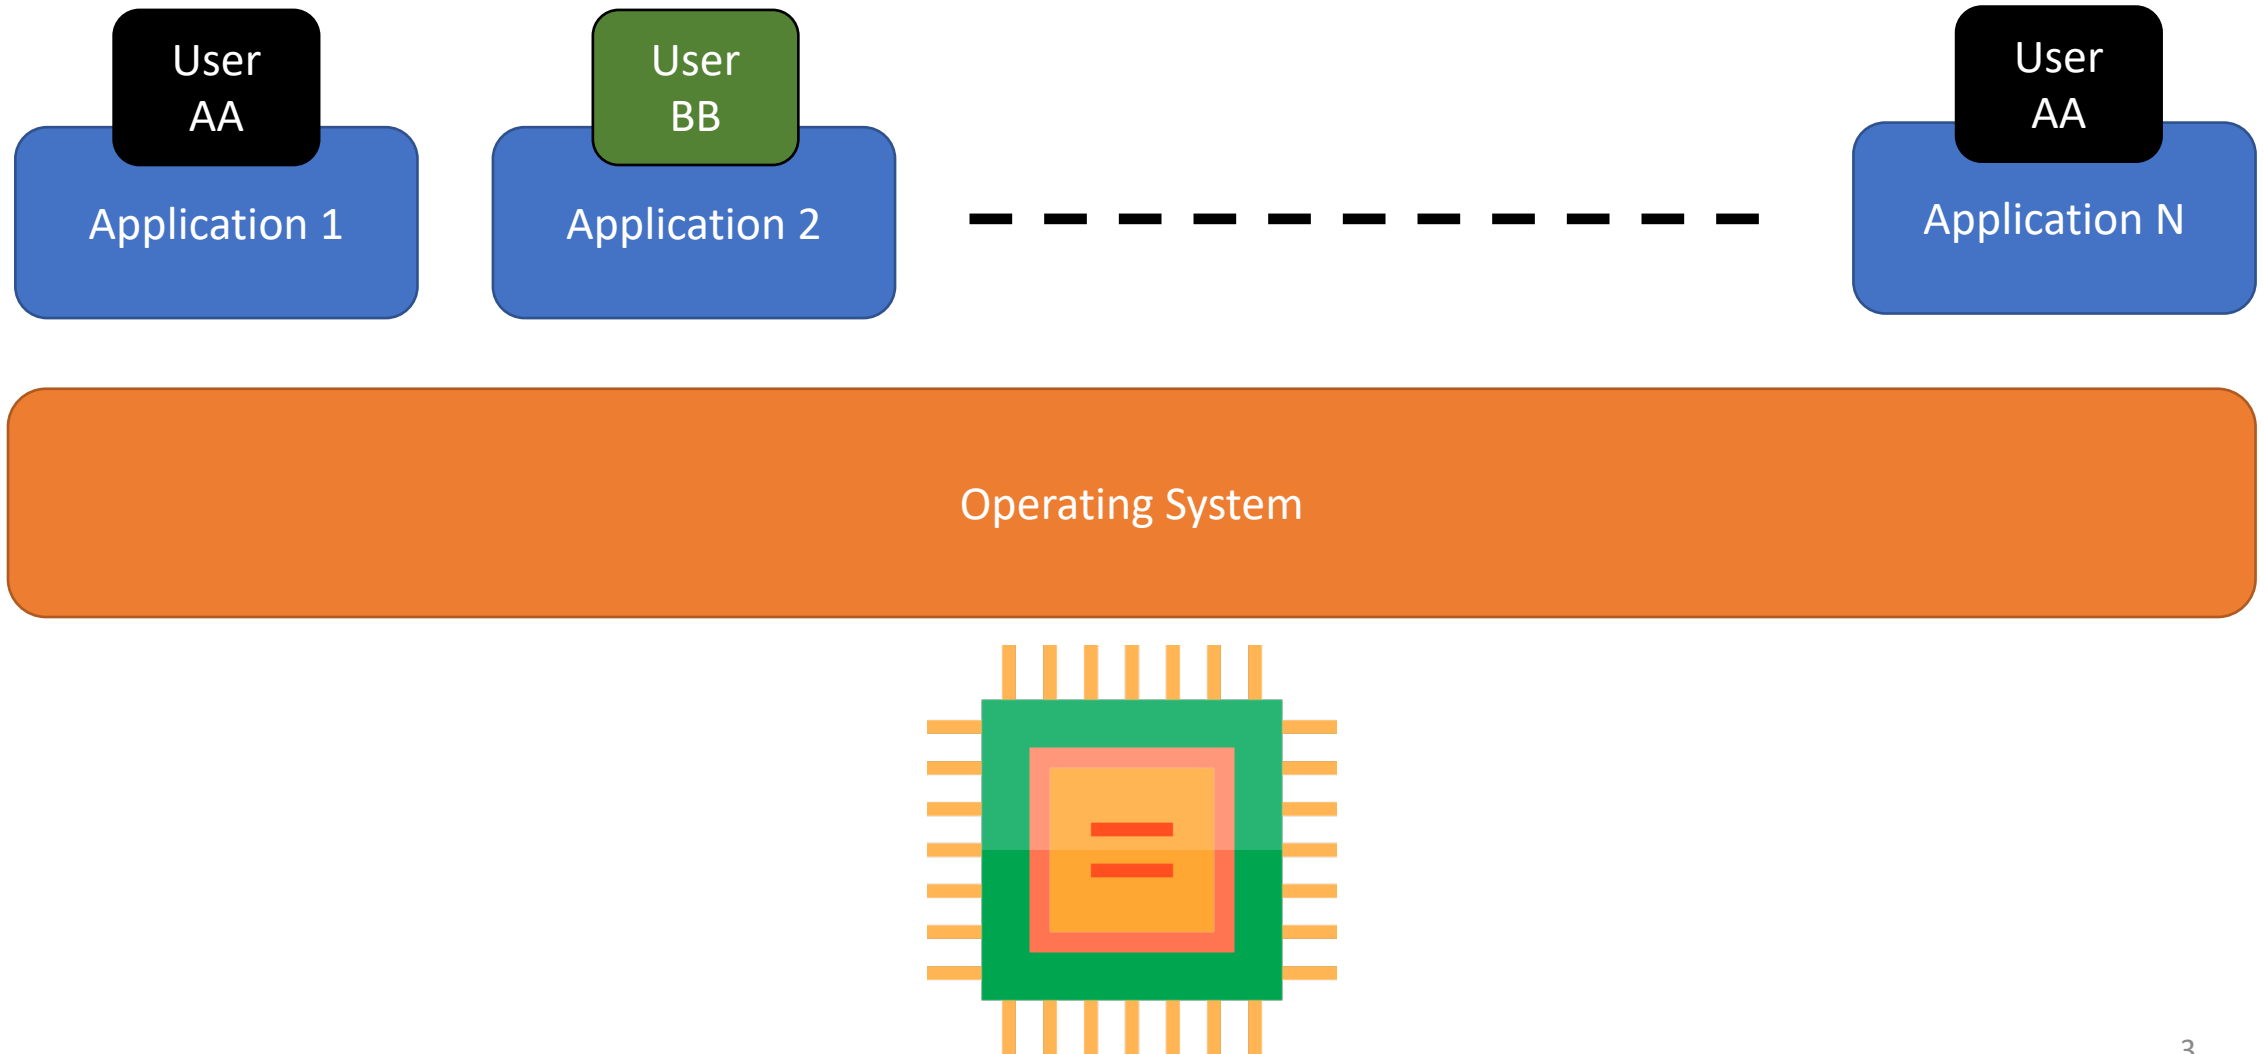
Lecture 12

Intro to Memory Management, Virtual Memory

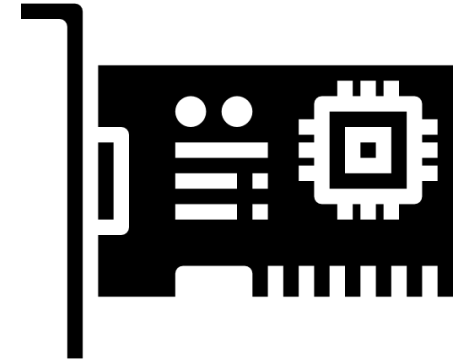
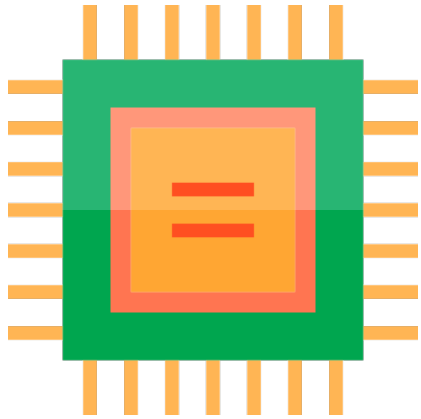
Logistics

- Lab 2 will be coming out this week
- Midterm Exam : Saturday, 18th March

Virtualizing the CPU



Major System Components



Let us talk about Memory Now

MANAGING DATA LOCALITY IN FUTURE
MEMORY HIERARCHIES USING A
HARDWARE SOFTWARE
CODESIGN APPROACH

by

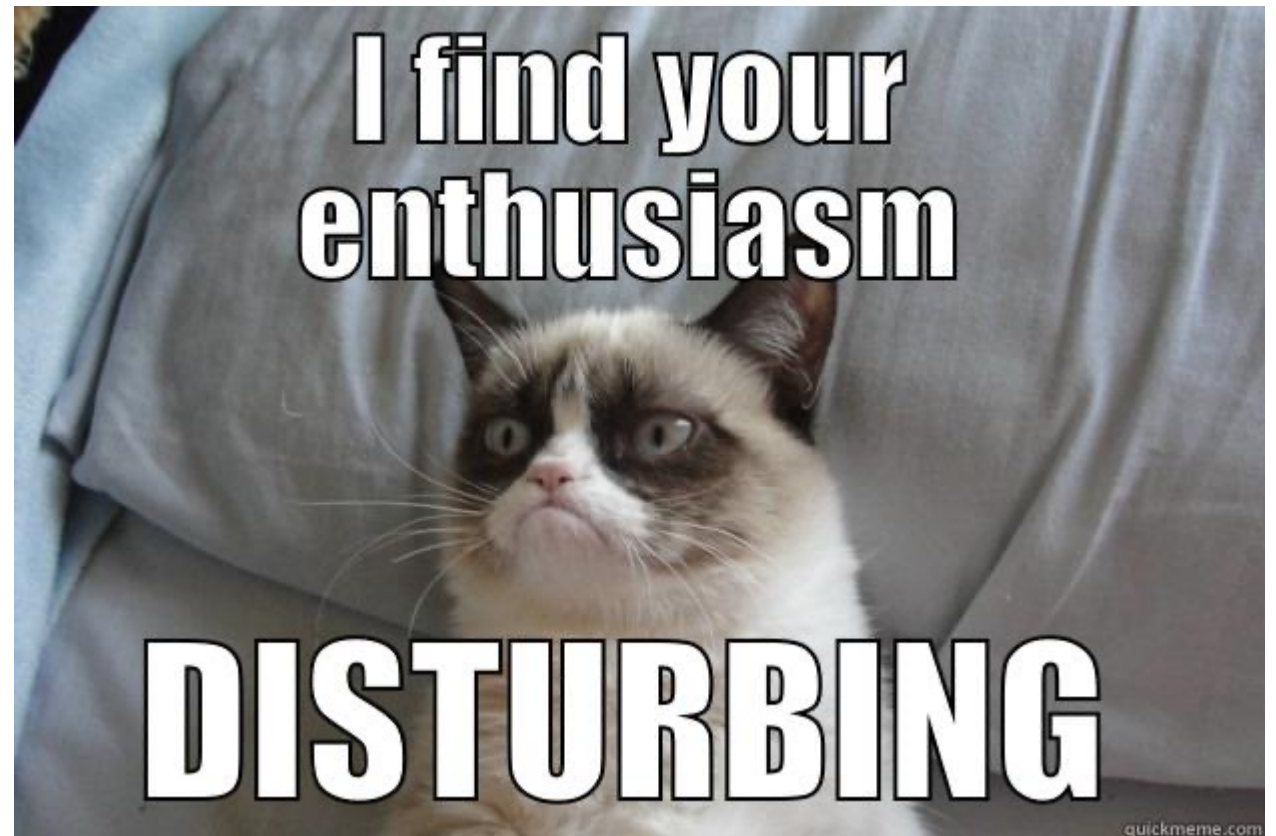
Manu Awasthi

A dissertation submitted to the faculty of
The University of Utah
in partial fulfillment of the requirements for the degree of

Doctor of Philosophy

in

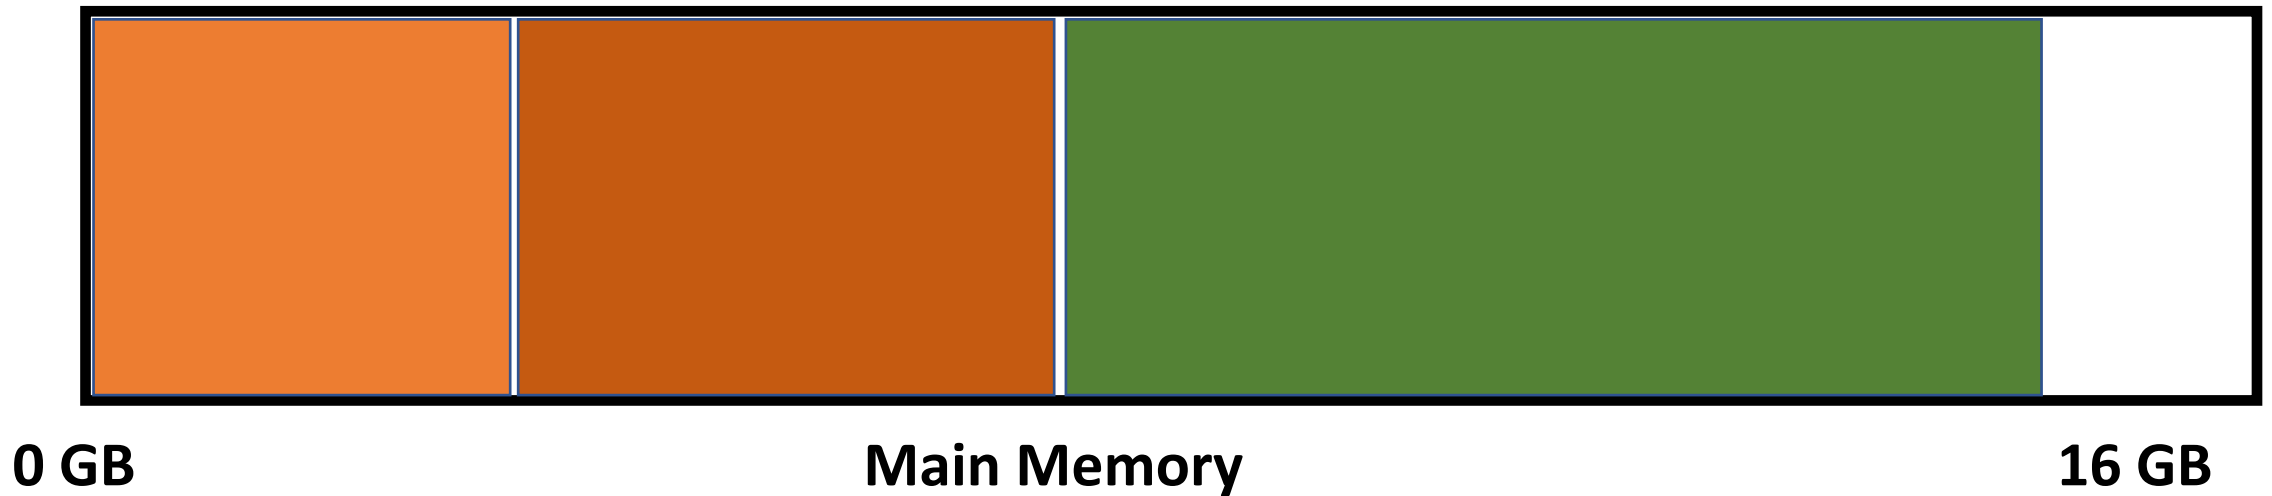
Computer Science



Spatio - Temporal Resource Multiplexing

- Resources can be multiplexed in two ways: **space** and **time**
- **Time Multiplexing:** sharing a resource by dividing up access to it over time.
 - Ride “sharing”
 - CPU scheduling on a single-core system
- **Space Multiplexing:** sharing a resource by dividing it into smaller pieces.
 - Shared apartments
 - Memory management

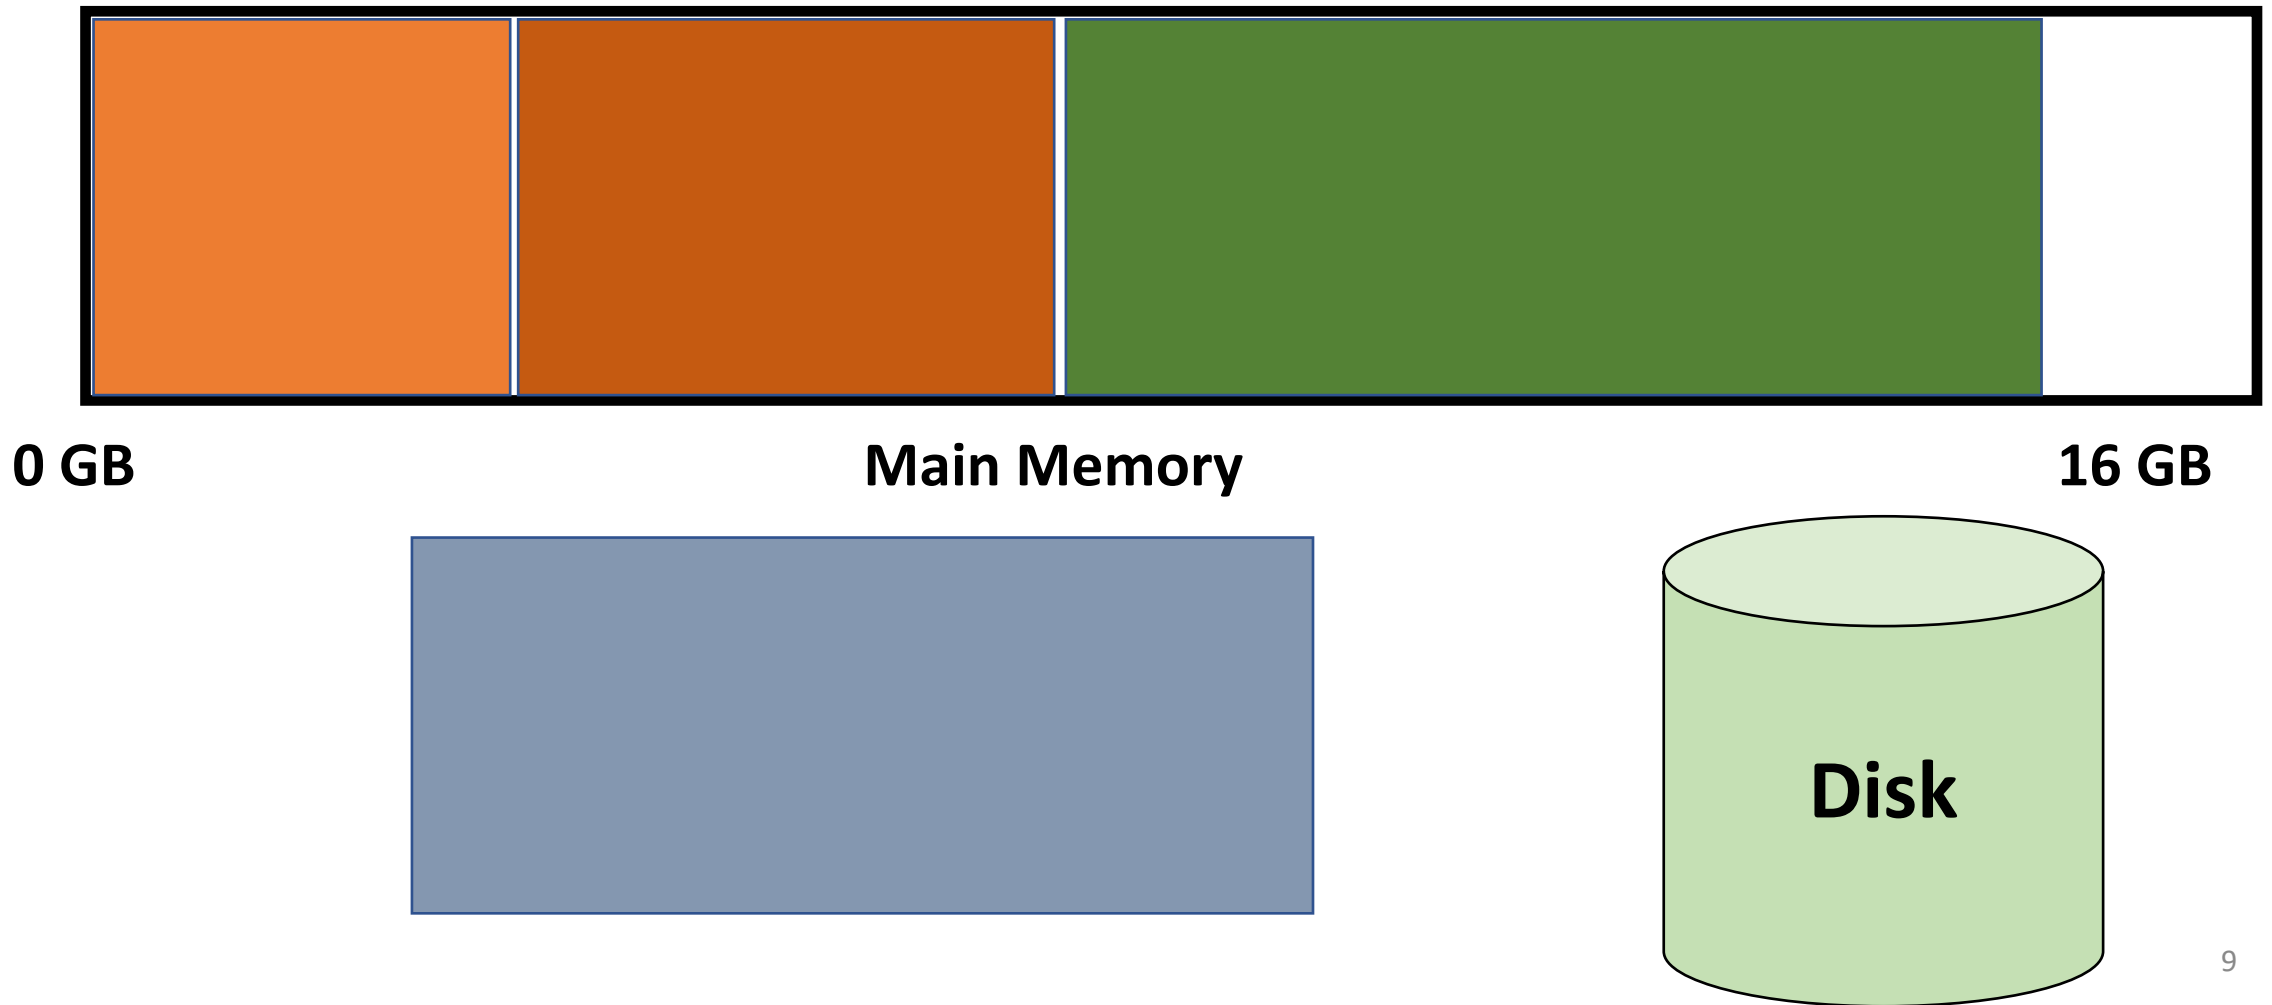
Chunk up Physical Memory



Issues

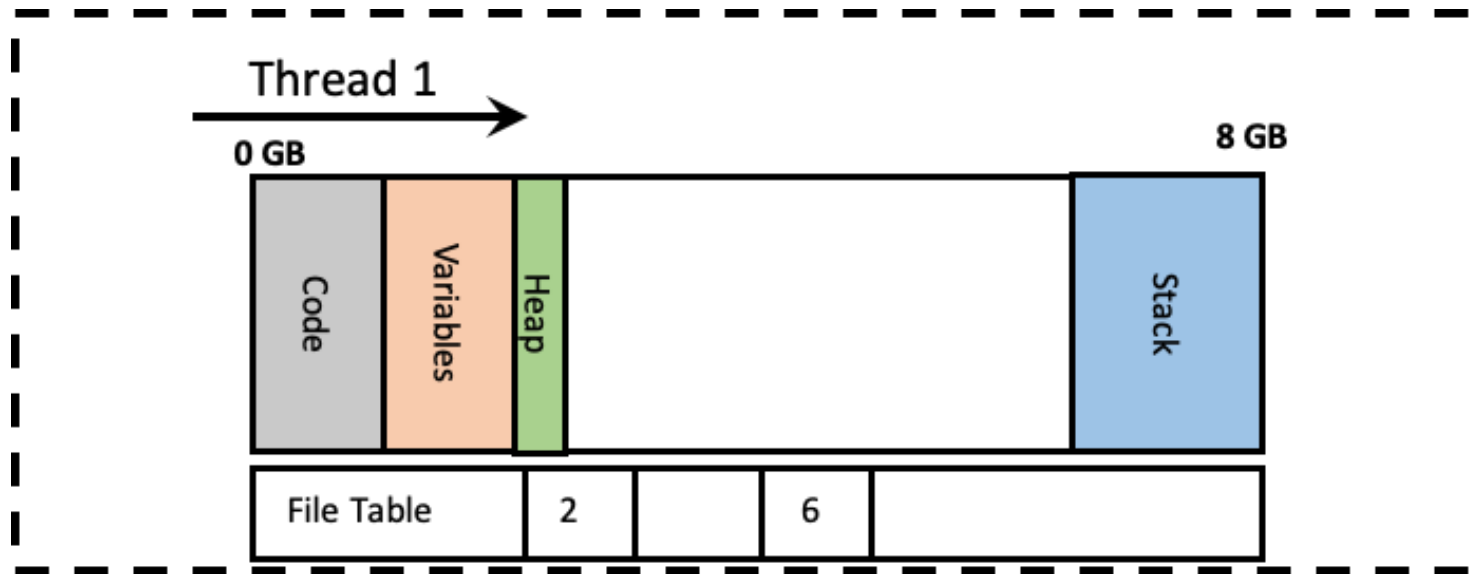
- Allocating Process some more memory at run time?
- Fragmentation
 - Internal
 - External

Capacity Issues



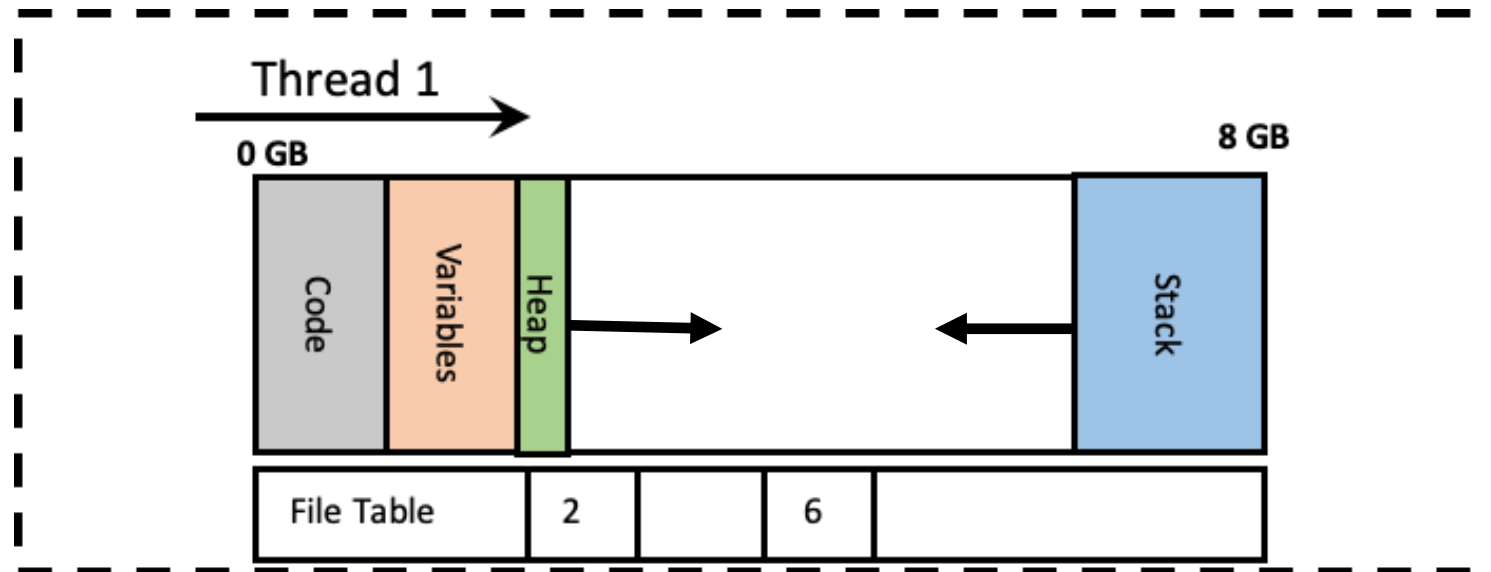
Process Address Spaces

- Processes are under illusion that they have access to **entire, contiguous** address space



How does this help?

- The uniformity of address spaces simplifies process **layout**:
 - "I always put my code and static variables at 0x10000."
 - "My heap always starts at 0x20000000 and grows **up**."
 - "My stack always starts at 0xFFFFFFFF and grows **down**."



Benefits

- Process layout is specified by Executable and Linker Format (ELF) file.
- Some layout is the function of **convention**.

```
cs304@cs304-devel:~/test$ readelf -l cpu_test | more

Elf file type is EXEC (Executable file)
Entry point 0x400a30
There are 6 program headers, starting at offset 64

Program Headers:
  Type           Offset             VirtAddr           PhysAddr
                 FileSiz             MemSiz             Flags   Align
LOAD             0x0000000000000000 0x0000000000400000 0x0000000000400000
                 0x000000000000b56b6 0x000000000000b56b6 R E     0x200000
LOAD             0x000000000000b6120 0x000000000006b6120 0x000000000006b6120
                 0x00000000000051b8 0x00000000000068e0 RW      0x200000
NOTE            0x0000000000000190 0x00000000000400190 0x00000000000400190
                 0x0000000000000044 0x0000000000000044 R        0x4
TLS             0x000000000000b6120 0x000000000006b6120 0x000000000006b6120
                 0x0000000000000020 0x0000000000000060 R        0x8
GNU_STACK       0x0000000000000000 0x0000000000000000 0x0000000000000000
                 0x0000000000000000 0x0000000000000000 RW       0x10
GNU_RELRO       0x000000000000b6120 0x000000000006b6120 0x000000000006b6120
                 0x0000000000002ee0 0x0000000000002ee0 R        0x1
```

Other Benefits : Relocatibility

```
int data[128];  
  
data[5] = 8; // Where is data[5]?  
  
result = bar(data[28]); // Where is bar?
```

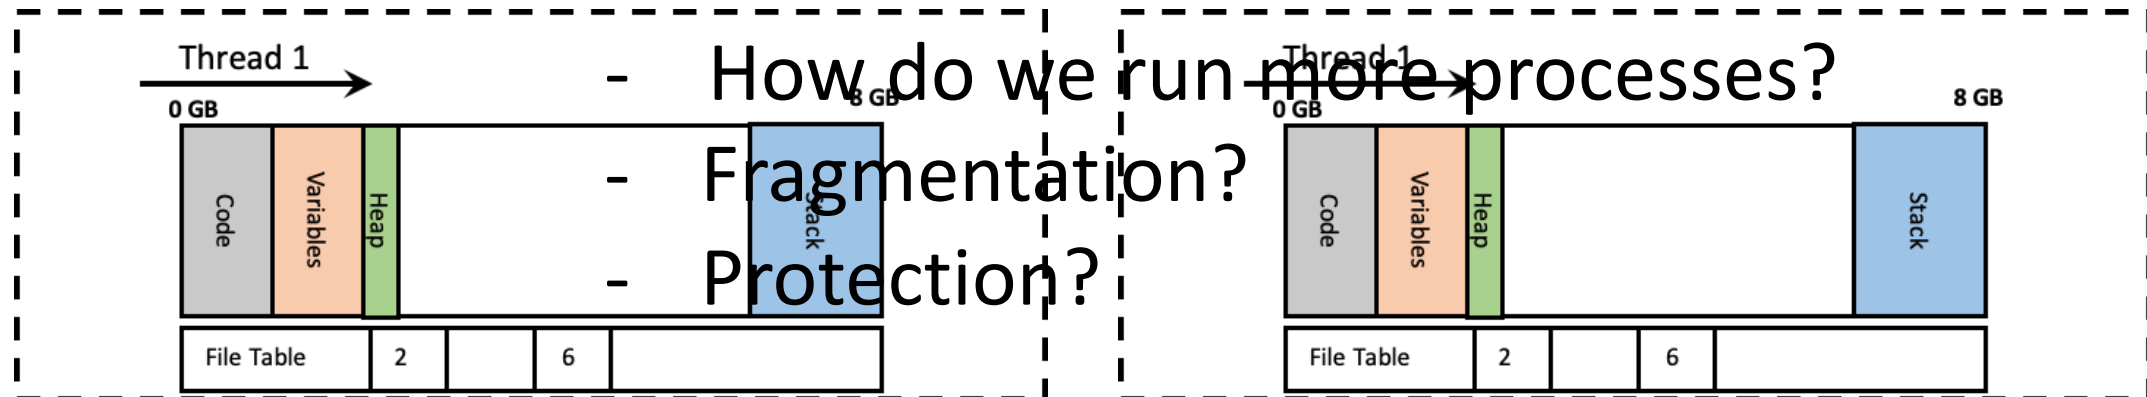
Implementation Issues



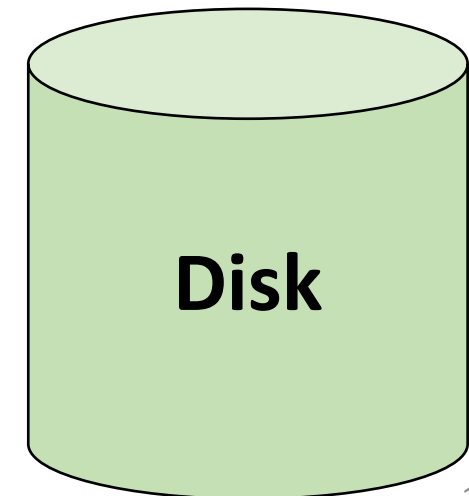
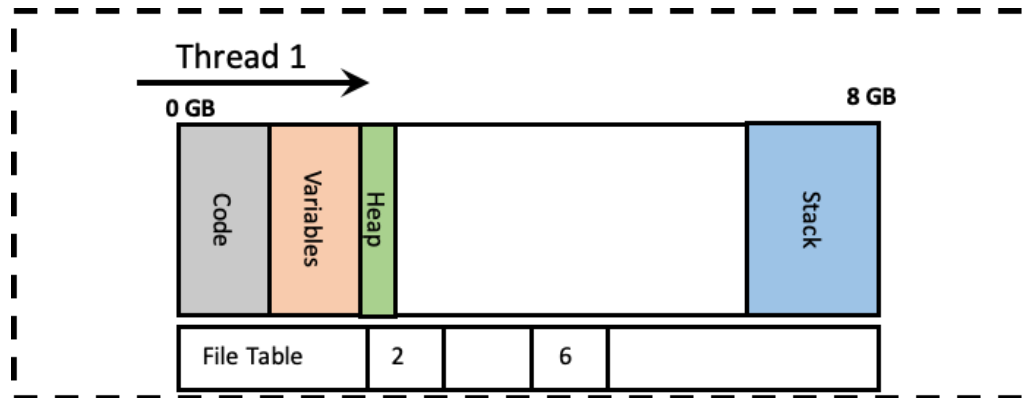
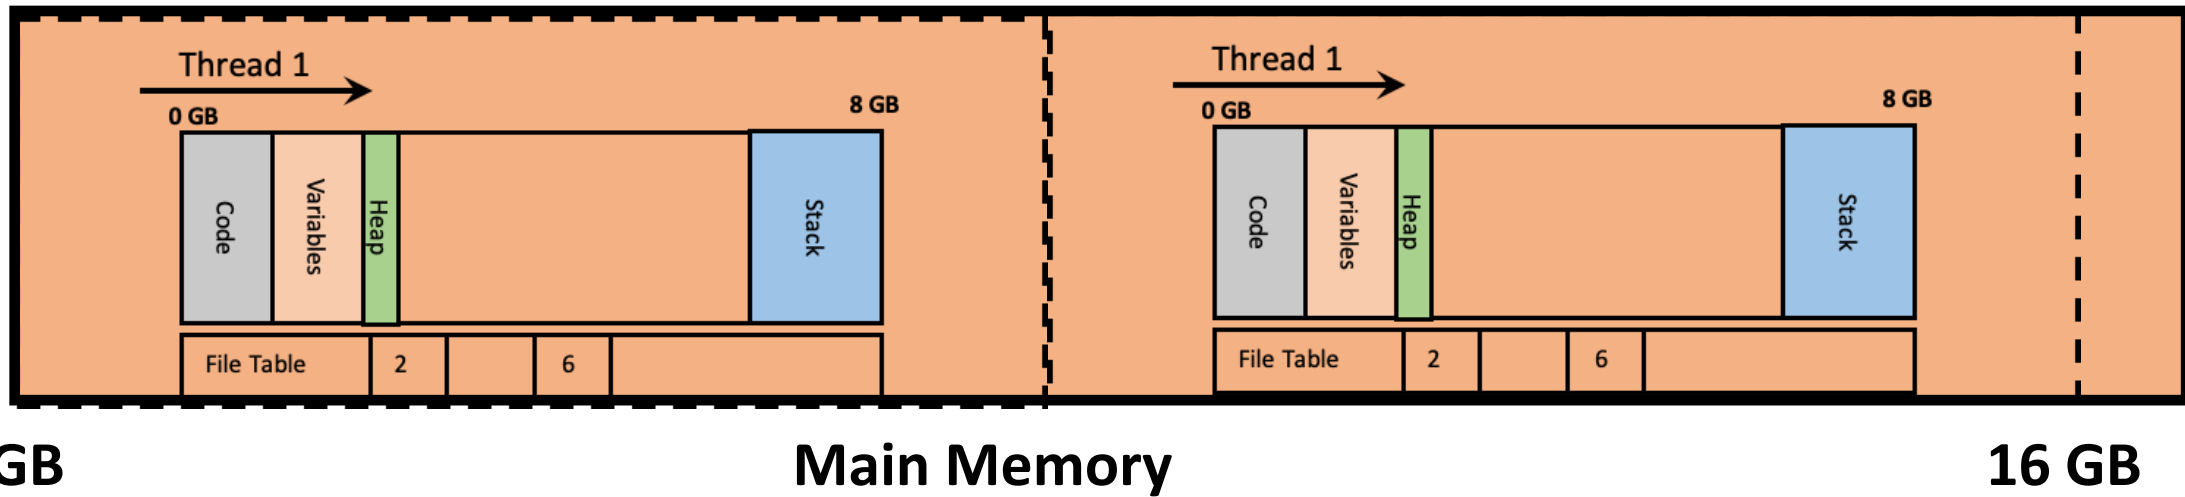
0 GB

Main Memory

16 GB

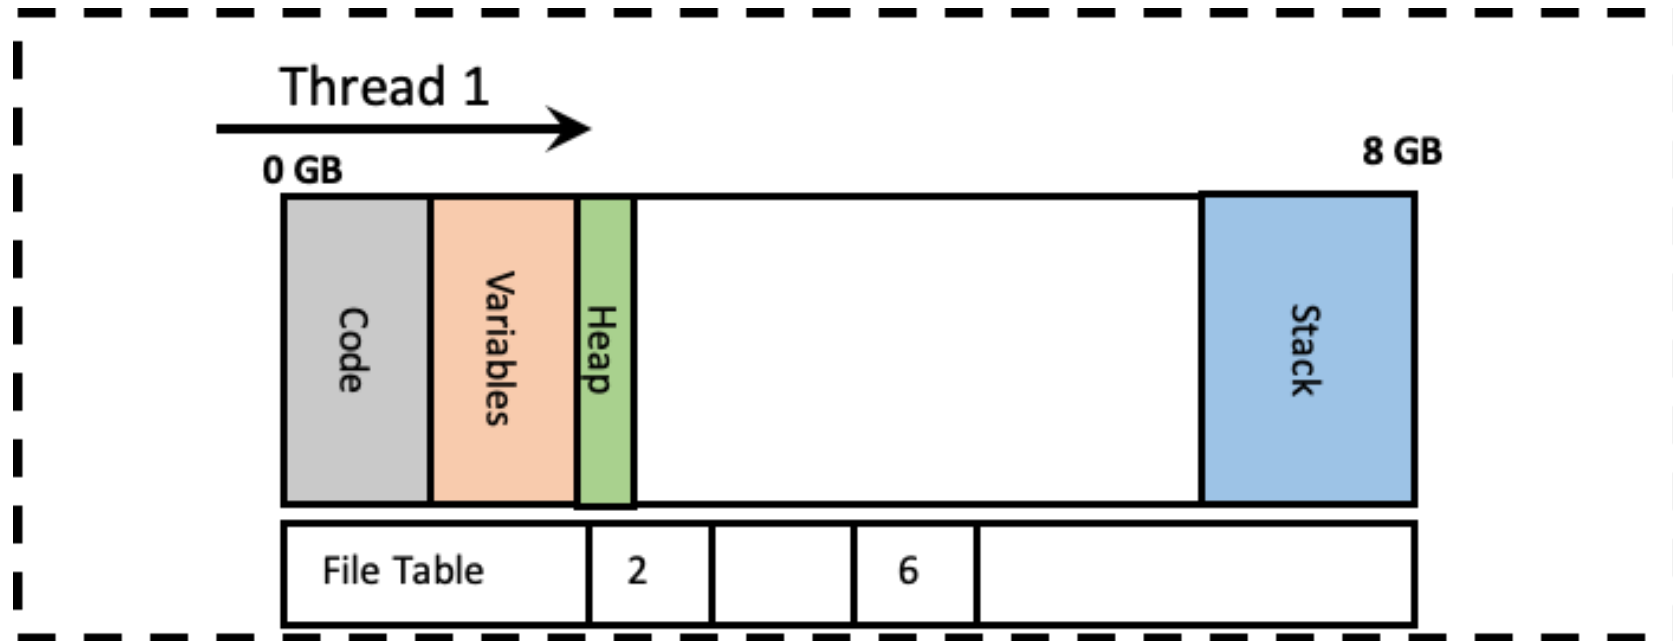


Solutions



Other Issues

- Process **might not need all** the space that it thinks it has.
- Managing the memory at a finer granularity might help



Solution : Add another level of indirection

- “All problems in computer science can be solved by another level of indirection” – David Wheeler

“Fundamental theorem of software engineering”



“Virtual” Address Spaces

- **Address translation:** 0x10000 to Process 1 is not the same as 0x10000 to Process 2 is not the same as...
- **Protection:** address spaces are intended to provide a *private* view of memory to each process.
- **Memory management:** together one or several processes may have **more address space** allocated than physical memory on the machine.