

CS 1217

Lecture 16 – Paging, Page Table Entries

Pages : Recap

- OS granularity of memory management
- 4KB typical size, other sizes also used
- Pages: 4KB sized allocations of memory (segments), where the **bound** is fixed
- Basic idea:
 - Divide Virtual **and** Physical Address Spaces into **Page** sized chunks
 - Establish one – to – one mapping between virtual and physical pages

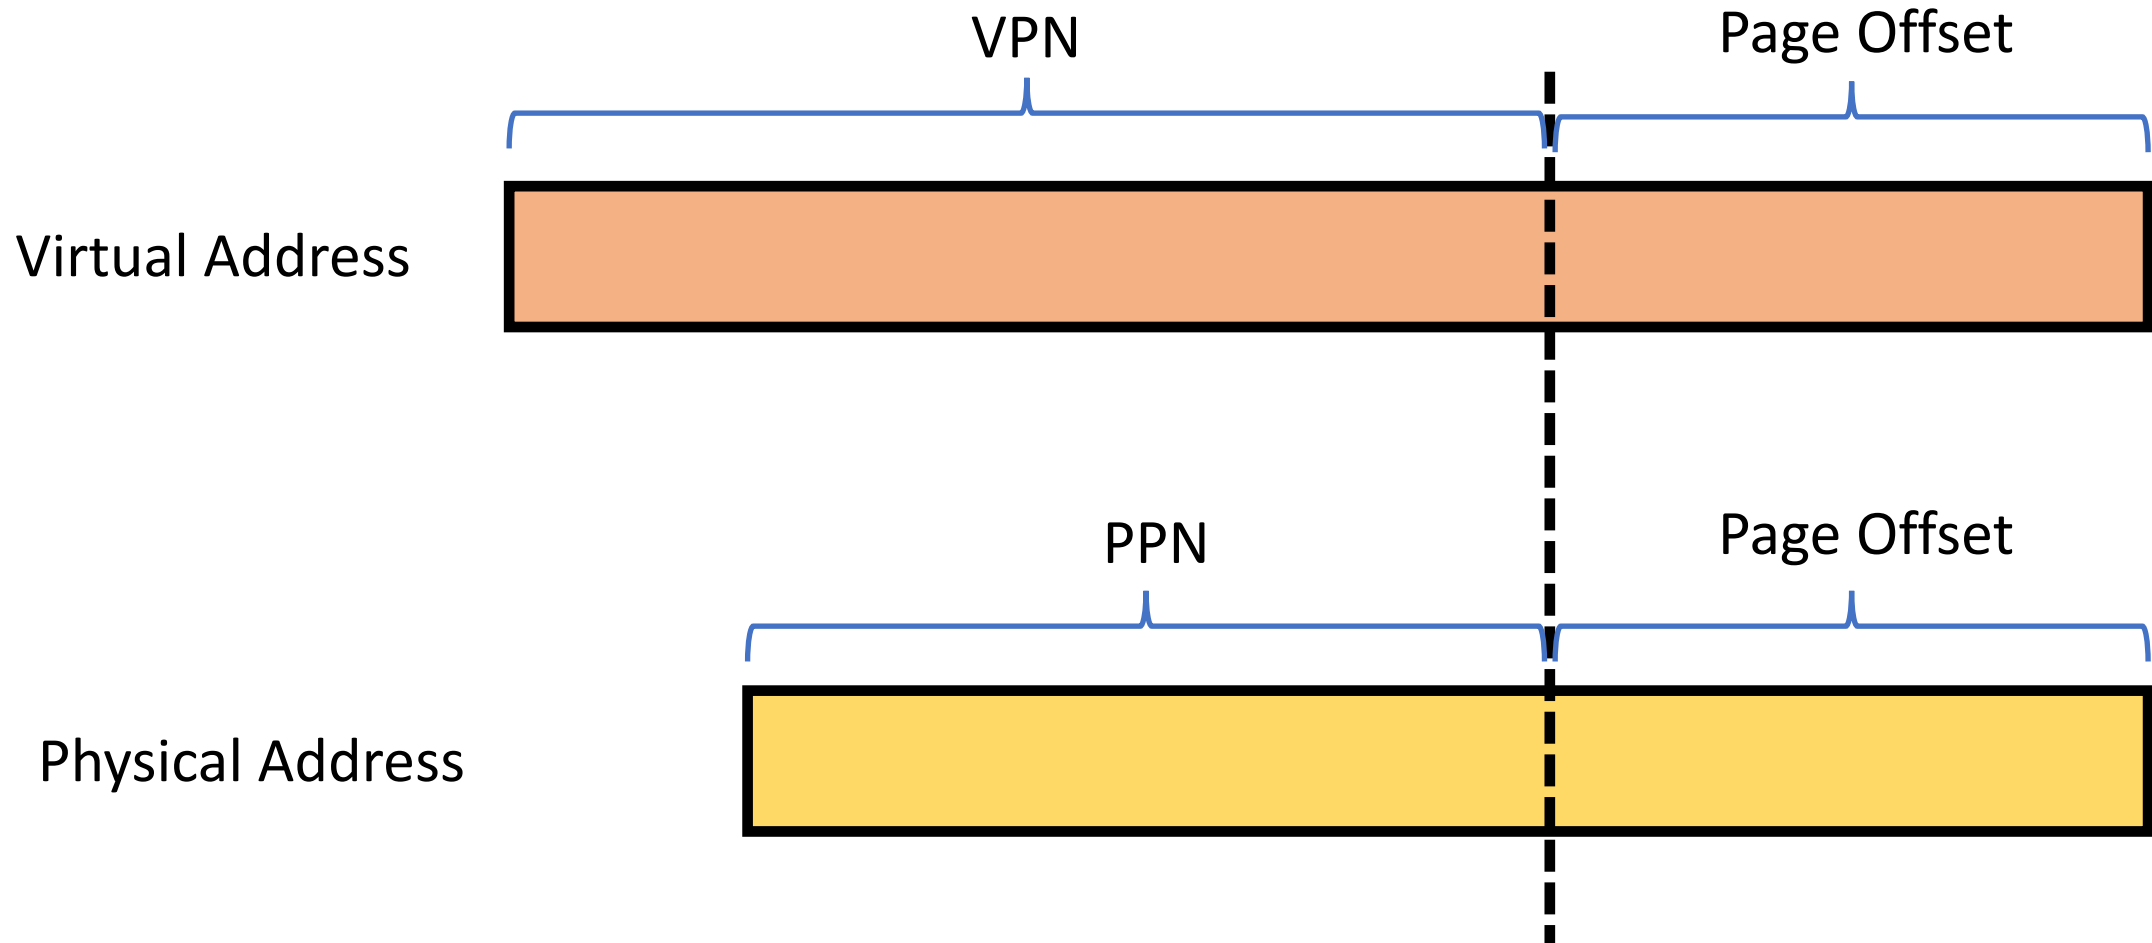
Address Translation using Pages

- Assuming 4K pages, 32-bit virtual memory address space
 - What is the size of the virtual address space?
 - How many bytes can be addressed per page?
 - How many bits are needed for the offset?
- **Check:** Get VPN, Check if a virtual to physical page translation exists
- **Translate:** Physical Address = Physical Page Number + Offset.



Virtual Address

Virtual Address to Physical Address



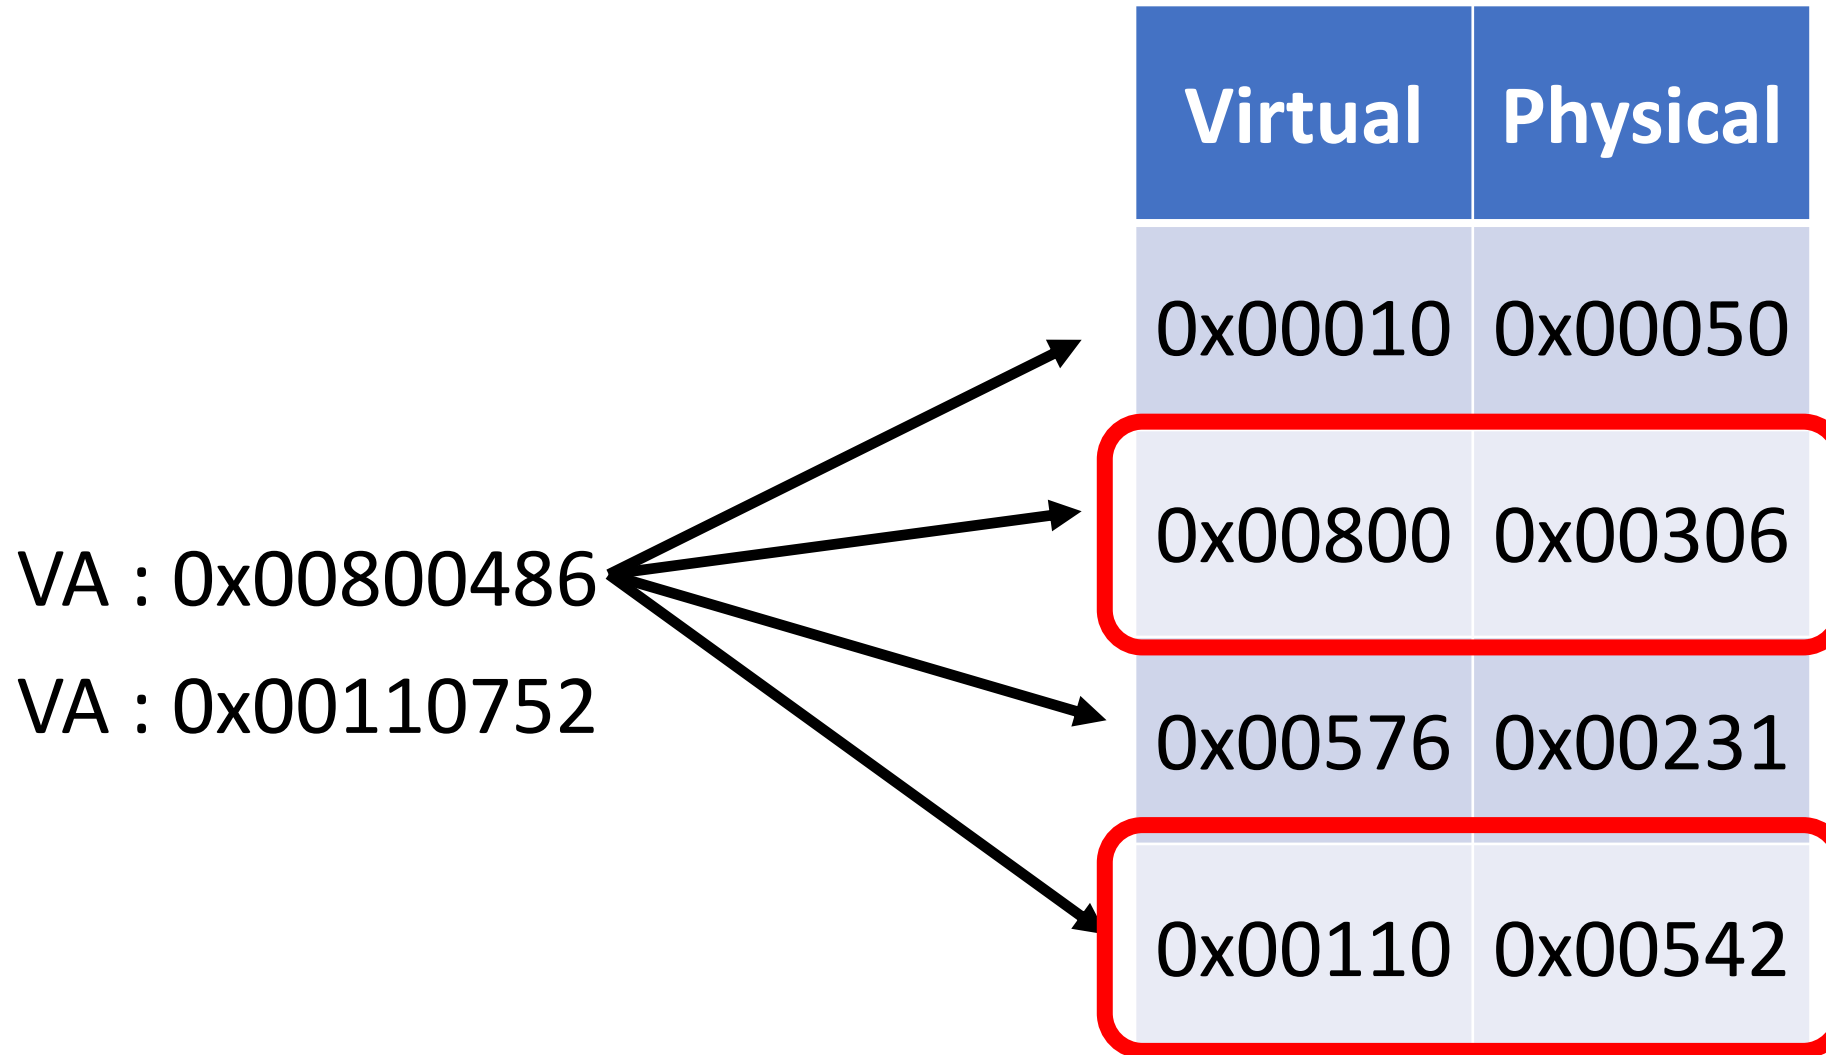
Address Translation using Pages

- **Address Translation**
- **Check:** Get VPN, Check if a virtual to physical page translation exists
- **Translate:** Physical Address = Physical Page Number + Offset.
- Assuming 4K pages, 32-bit virtual memory address space
 - What is the size of the virtual address space?
 - How many bytes can be addressed per page?
 - How many bits are needed for the offset?

Virtual Address



Role of TLBs now




Pros and Cons

- Access permissions?
 - Can now be associated on a per-page basis
- External fragmentation?
 - None: All allocations are multiples of Page Size (e.g. 4KB)
- Internal fragmentation?
 - Maybe a little, but much reduced as compared to segmentation
- Con: requires per-page hardware translation, need h/w help
- Con: requires per-page operating system state (page tables)

Page based Address Translations

- TLBs **cache** virtual to physical translations
- Where are the translations really located?
 - i.e., the ones that are **not cached** in the TLB
- **Page Table Entries (PTEs)**
 - a single-entry storing information about a single virtual page used by a single process
 - Information stored?



Page Table Entries

- Page Table Entry (PTE) = A single entry, storing information about a virtual page
- Can contain a lot of information
 - Location – which physical page does this virtual page map to
 - Permissions – Read/Write/Execute
 - Present – whether this page is present in memory
 - Dirty – has this page been written to recently?


The Brave New (World) Translation Path

- *Process*: Store to address 0x10000
- *MMU*: I don't know which physical address 0x10000 maps to
- **Exception!**
- *Kernel*: Need to look up PTE to which 0x10000 belongs
- What should be the properties of this lookup?
 - Fast (Time efficiency)
 - Compact (Space efficiency)

Page Tables

- In essence, a data structure, which
- Maps : a **virtual page number** to a **page table entry**
- Each process has its own page table
 - Why?

Page Table Organization

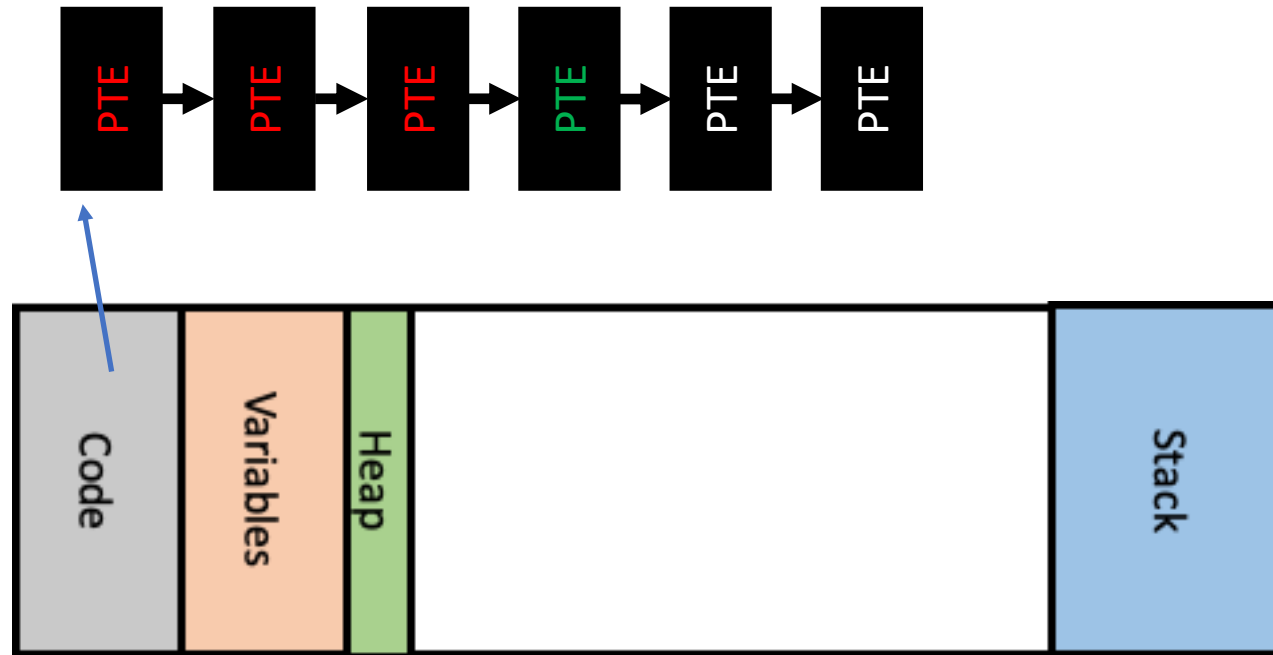


Flat Page Tables

- Time Complexity
 - $O(1)$
 - Fast!
- Space Complexity
 - 4 MB per process
 - Gets expensive really fast with number of processes

Linked List Page Table

- Maintain LL of **valid** PTEs, to be searched on each lookup



Linked List Page Table


- Time Complexity
 - $O(N)$ for N allocated virtual pages
- Space Complexity
 - $4B * N$ per process

Multi Level Page Tables

- Insight: Utilize the sparsity of virtual address space
- Key Idea: Break VPN into multiple parts, each used as an index at a separate level of the tree.
- Example: with 4K pages VPN is **20** bits. Use **top 10** bits as index into **top-level** page table, **bottom 10 bits** as index into second-level page table



Multi Level Page Tables



Multi Level Page Tables

- Time Complexity:
 - $O(c)$. Constant number of look ups per translation depending on tree depth
-
- Space Complexity:
 - Depends on sparsity of address space, but better than flat.