

# CS 1217

Swapping Wrap; Page Replacement Policies; Free Space Management

# Logistics

- Lab 3 is out today; might have one more assignment on Synchronization
- 1 make-up lecture during reading week, followed by final exam in same week

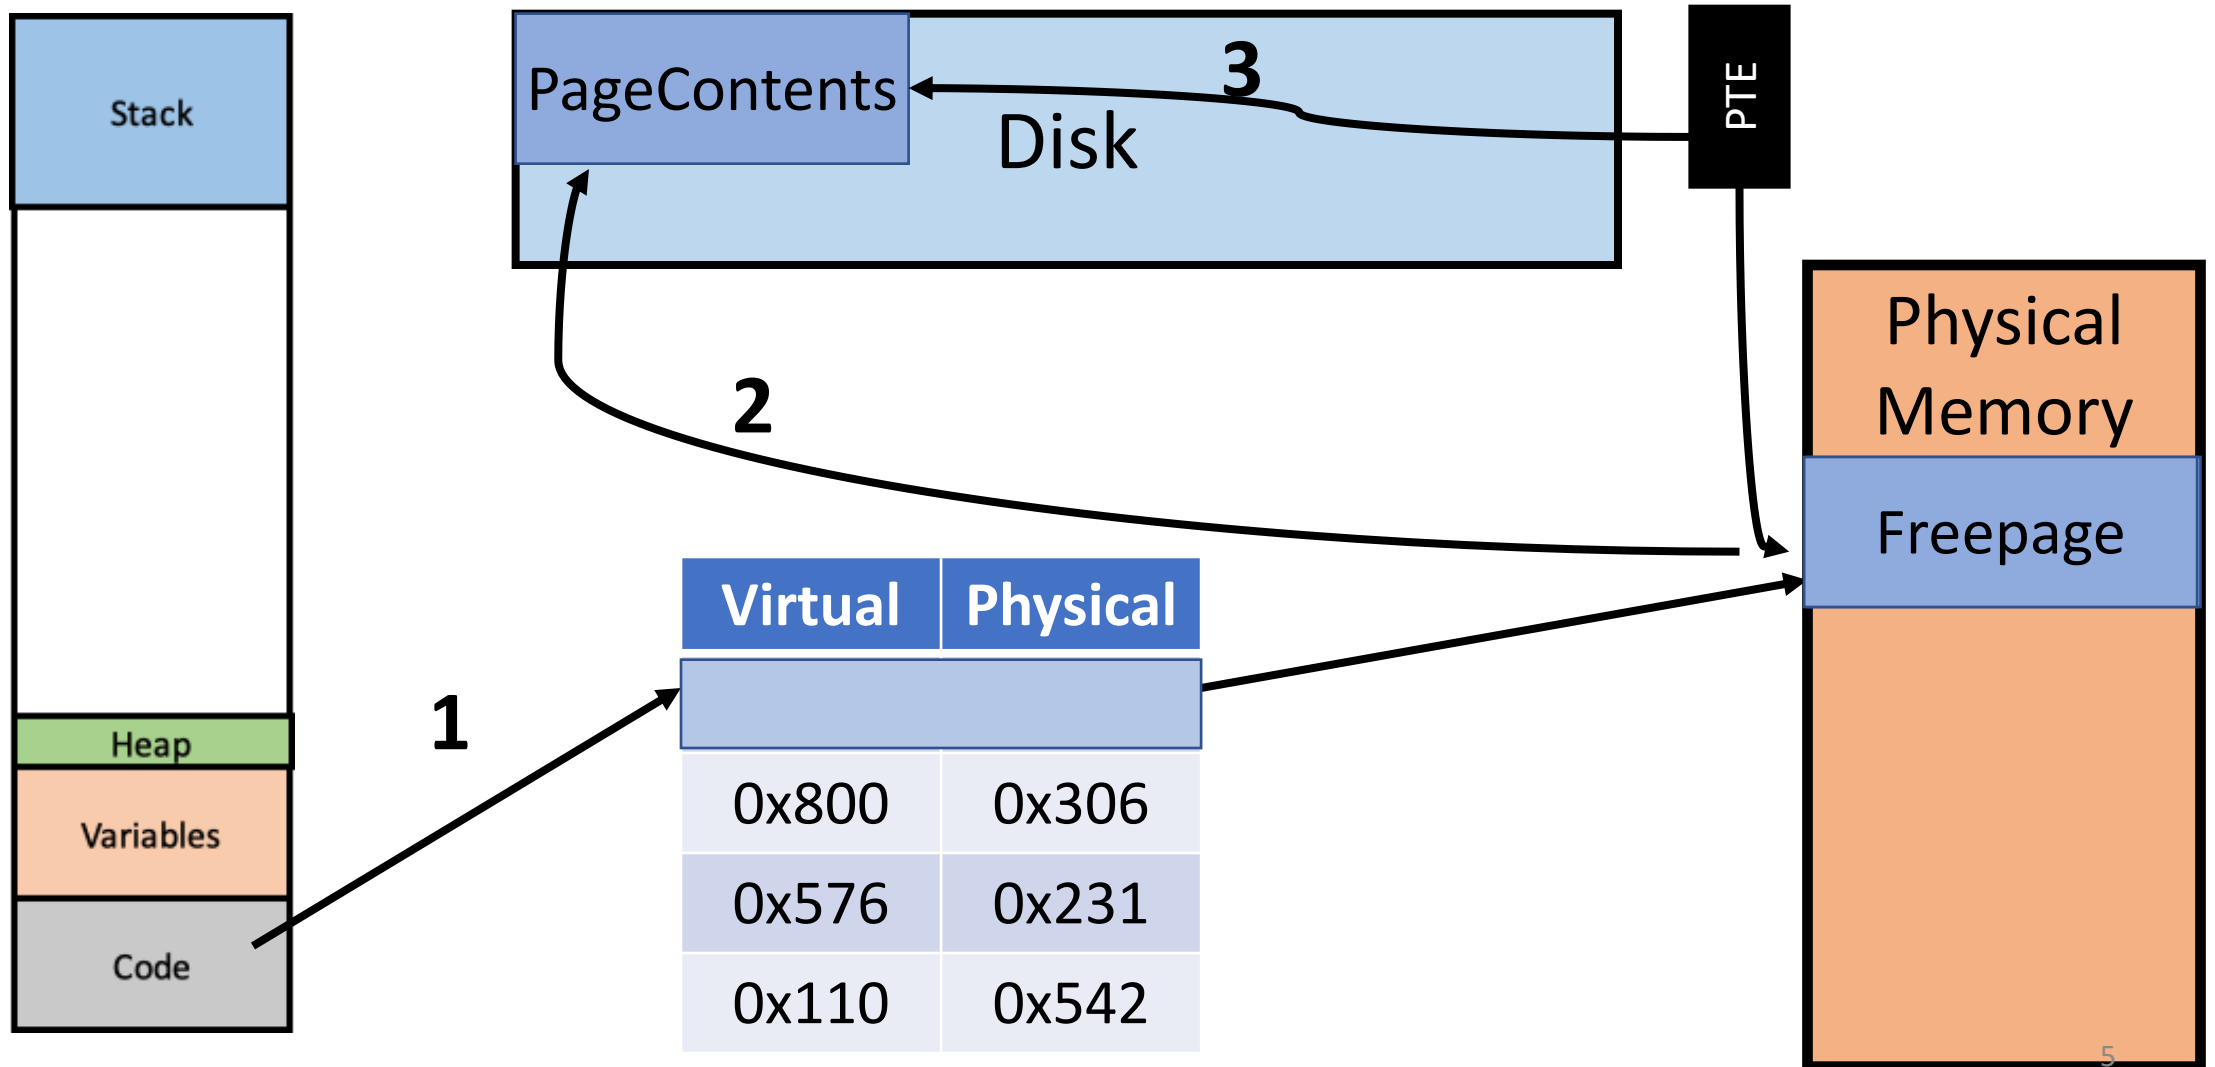
# Recap: Swapping

- Swap Space: The on-disk space that operating systems typically place data stored in memory in order to *borrow* memory.
- Swapping: Process of moving (swapping) data back and forth between memory and disk
- Goal: make it **feel** like the system has memory that is as *large* as the size of the disk and as *fast* as actual RAM

# Swap Out Process

- **Remove** the translation from the TLB, if it exists
- **Copy** the contents of the page to disk.
- **Update** the page table entry to indicate that the page is on disk

# Swap Out Process



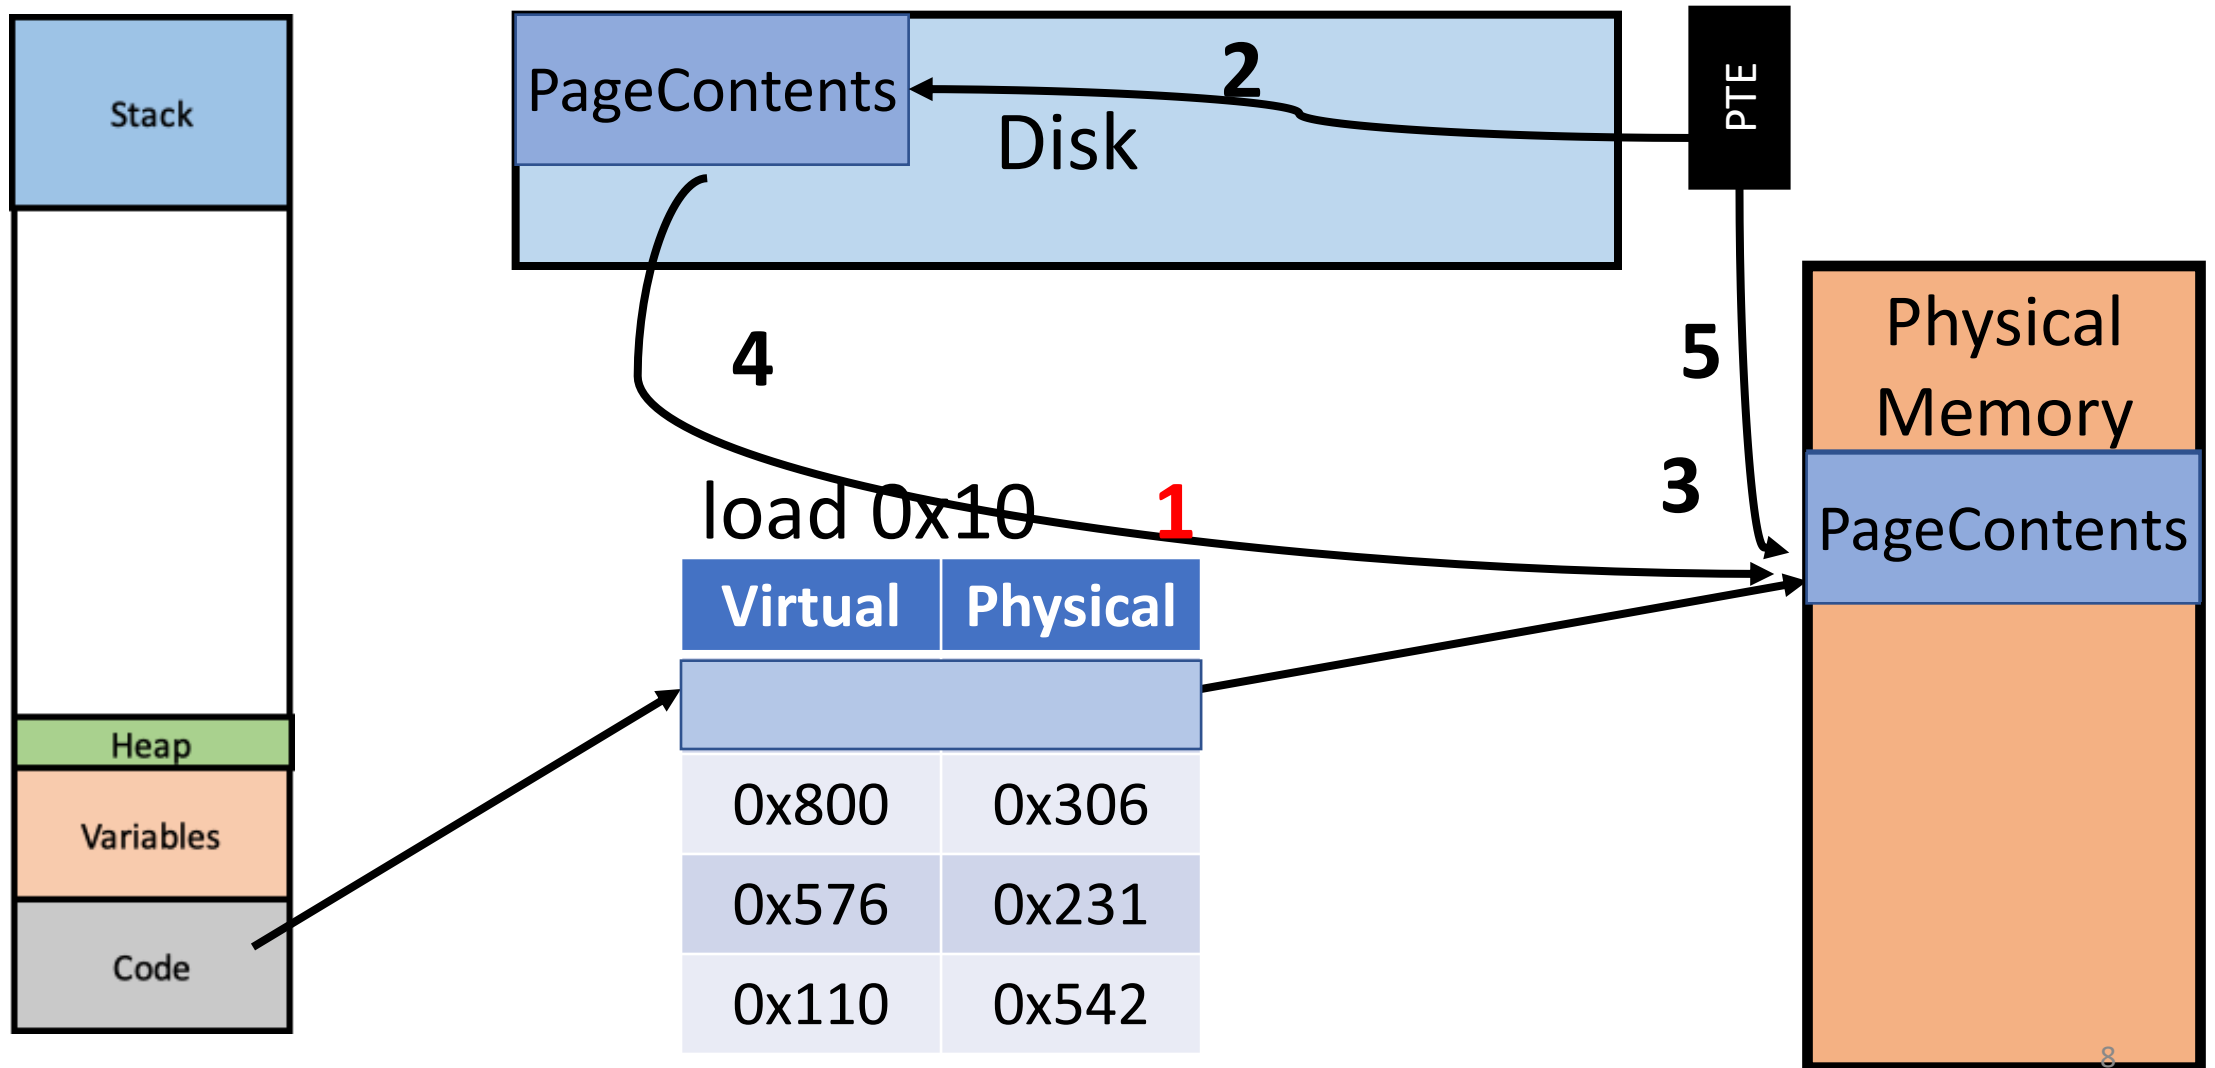
# Swap Out Speed

- **Remove** the translation from the TLB
  - Fast or Slow?
- **Copy** the contents of the page to disk
  - Fast or Slow?
- **Update** the page table entry to indicate that the page is on disk
  - Fast or Slow?

# Swap In

- When should swap in be done?
  - When the virtual address is used by the process!
- To translate an address for a process that has been swapped out:
  - **Stall** the instruction till contents of page are retrieved
  - **Allocate** a page in memory to hold the new page contents.
  - **Locate** the page on disk using the page table entry.
  - **Copy** the contents of the page from disk.
  - **Update** the page table entry to indicate that the page is in memory
  - **Load** the TLB
  - **Restart** the instruction using the retrieved virtual address.

# Swap In Process





# On Demand Paging

- Until an instruction on a code page is **executed**, or
  - a **read** or **write** occurs to a data or heap page
  - the kernel does **not** load the contents of that page into memory!
- 
- WHY?
  - A **lot** of code is never executed and some global variables are never used. **Why waste memory?**
  - Map virtual addresses to physical addresses “on demand”

# On Demand Paging

- What happens the first time a process executes an instruction from a **new code page**?
  - That page contents are loaded from disk and the instruction is restarted
- What happens the first time a does a load or store to an uninitialized heap, stack or data page?
  - The kernel allocates a **new page** filled with appropriate information and the instruction is restarted.

# Page Eviction: What can go wrong?

***Thrashing** is a colloquialism normally used to describe a computer whose virtual memory subsystem is in a constant state of paging, rapidly exchanging data in memory for data on disk, to the exclusion of most application-level processing. This causes the performance of the computer to degrade or collapse.*

- Wikipedia