

CS1217 - Spring 2023 - Homework 3

Gautam Ahuja, Nistha Singh

Note: All screenshots are on Gautam Ahuja's machine.

Boot xv6

Done by: Nistha Singh

To boot the xv6 operating system, we follow the procedure explained in the assignment.

We run the `make` command on the terminal.

```
gautam-ahuja@LAPTOP-FV7627LB:~/.../cs1217-assignment-3-julius-stabs-back$ make
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -fno-p
ic -O -nostdinc -I. -c bootmain.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -fno-p
ic -nostdinc -I. -c bootasm.S
ld -m elf_i386 -N -e start -Ttext 0x7C00 -o bootblock.o bootasm.o bootmain.o
objdump -S bootblock.o > bootblock.asm
objcopy -S -O binary -j .text bootblock.o bootblock
./sign.pl bootblock
boot block is 451 bytes (max 510)
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o bio.o bio.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o console.o console.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o exec.o exec.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o file.o file.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o fs.o fs.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o ide.o ide.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o ioapic.o ioapic.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o kalloc.o kalloc.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o kbd.o kbd.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o lapic.o lapic.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o log.o log.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o main.o main.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o mp.o mp.c
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie -c -
o picirq.o picirq.c
```

We run the `nm kernel | grep _start` to access the starting address of the kernel.

```
gautam-ahuja@LAPTOP-FV7627LB:~/.../cs1217-assignment-3-julius-stabs-back$ nm
kernel | grep _start
8010a48c D _binary_entryother_start
8010a460 D _binary_initcode_start
0010000c T _start
gautam-ahuja@LAPTOP-FV7627LB:~/.../cs1217-assignment-3-julius-stabs-back$
```

Now we run two commands side by side: `make qemu-nox-gdb` and `gdb` on different terminals. This will link the gdb to kernel.

```

gautam-ahuja@LAPTOP-FV76  X  +  v
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o userstests.o userstests.c
ld -m elf_i386 -N -e main -Ttext 0 -o _userstests userstests.o ulib.o usys.o
printf.o umalloc.o
objdump -S _userstests > userstests.asm
objdump -t _userstests | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$/d' > userstest
s.sym
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o wc.o wc.c
ld -m elf_i386 -N -e main -Ttext 0 -o _wc wc.o ulib.o usys.o printf.o uma
lloc.o
objdump -S _wc > wc.asm
objdump -t _wc | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$/d' > wc.sym
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o zombie.o zombie.c
ld -m elf_i386 -N -e main -Ttext 0 -o _zombie zombie.o ulib.o usys.o prin
tf.o umalloc.o
objdump -S _zombie > zombie.asm
objdump -t _zombie | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$/d' > zombie.sym
gcc -fno-pic -static -fno-builtin -fno-strict-aliasing -O2 -Wall -MD -ggdb -
m32 -Werror -fno-omit-frame-pointer -fno-stack-protector -fno-pie -no-pie
-c -o date.o date.c
ld -m elf_i386 -N -e main -Ttext 0 -o _date date.o ulib.o usys.o printf.o
umalloc.o
objdump -S _date > date.asm
objdump -t _date | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$/d' > date.sym
./mkfs fs.img README_cat_echo_forktest_grep_init_kill_ln_ls_mkdir _
rm _sh_stressfs _userstests _wc _zombie _date
nmeta 59 (boot, super, log blocks 30 inode blocks 26, bitmap blocks 1) block
s 941 total 1000
ballocc: first 673 blocks have been allocated
ballocc: write bitmap block at sector 58
sed 's/localhost:1234/localhost:26000/' < .gdbinit.tmpl > .gdbinit
*** Now run 'gdb'.
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw
-drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512 -S -gdb tc
p::26000

gautam-ahuja@LAPTOP-FV7627LB:~/.../cs1217-assignment-3-julius-stabs-back$ gd
b
GNU gdb (Ubuntu 12.1-0ubuntu1~22.04) 12.1
Copyright (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html
>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word".
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
The target architecture is set to "i8086".
[f000:ffff] 0xfffff0: jmp $0x3630,$0xf000e05b
0x0000ffff in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configu
ration
of GDB. Attempting to continue with the default i8086 settings.

(gdb) |

```

Now, we set the break point `br * 0x010000c` and continue to boot up the xv6 operating system.

```

gautam-ahuja@LAPTOP-FV76  X  +  v
SeaBIOS (version 1.15.0-1)

iPXE (https://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8B4A0+1FECB4A0 CA00

Booting from Hard Disk..xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap8
init: starting sh
$ |

gautam-ahuja@LAPTOP-FV7627LB:~/.../cs1217-assignment-3-julius-stabs-back$ gd
b
GNU gdb (Ubuntu 12.1-0ubuntu1~22.04) 12.1
Copyright (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html
>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word".
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
The target architecture is set to "i8086".
[f000:ffff] 0xfffff0: jmp $0x3630,$0xf000e05b
0x0000ffff in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configu
ration
of GDB. Attempting to continue with the default i8086 settings.

(gdb) br * 0x010000c
Breakpoint 1 at 0x10000c
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x10000c: mov %cr4,%eax

Thread 1 hit Breakpoint 1, 0x010000c in ?? ()
(gdb) c
Continuing.
|

```

We have now booted up the xv6 OS.

0. Exercise 0

Done by: Nistha Singh

info reg gives the content of all the registers while running.

```
Thread 1 hit Breakpoint 1, 0x0010000c in ?? (C)
(gdb) info reg
(gdb) info reg
eax            0x0            0
ecx            0x0            0
edx            0x1f0          496
ebx            0x10094        65684
esp            0x7bdc         0x7bdc
ebp            0x7bf8         0x7bf8
esi            0x10094        65684
edi            0x0            0
eip            0x10000c        0x10000c
eflags         0x46          [ IOPL=0 ZF PF ]
cs             0x8            8
ss             0x10           16
ds             0x10           16
es             0x10           16
fs             0x0            0
gs             0x0            0
fs_base        0x0            0
gs_base        0x0            0
k_gs_base      0x0            0
cr0            0x11          [ ET PE ]
cr2            0x0            0
cr3            0x0            0
cr4            0x0            0
cr8            0x0            0
efer           0x0            0
xmm0           {v4_float = {0x0, 0x0, 0x0, 0x0}, v2_double = {0x0, 0x0}, v16_int8 = {0x0 <repeats 16 times>}, v8_int16 = {0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0},
v4_int32 = {0x0, 0x0, 0x0, 0x0}, v2_int64 = {0x0, 0x0}, uint128 = 0x0}
xmm1           {v4_float = {0x0, 0x0, 0x0, 0x0}, v2_double = {0x0, 0x0}, v16_int8 = {0x0 <repeats 16 times>}, v8_int16 = {0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0},
v4_int32 = {0x0, 0x0, 0x0, 0x0}, v2_int64 = {0x0, 0x0}, uint128 = 0x0}
xmm2           {v4_float = {0x0, 0x0, 0x0, 0x0}, v2_double = {0x0, 0x0}, v16_int8 = {0x0 <repeats 16 times>}, v8_int16 = {0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0},
v4_int32 = {0x0, 0x0, 0x0, 0x0}, v2_int64 = {0x0, 0x0}, uint128 = 0x0}
xmm3           {v4_float = {0x0, 0x0, 0x0, 0x0}, v2_double = {0x0, 0x0}, v16_int8 = {0x0 <repeats 16 times>}, v8_int16 = {0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0},
v4_int32 = {0x0, 0x0, 0x0, 0x0}, v2_int64 = {0x0, 0x0}, uint128 = 0x0}
xmm4           {v4_float = {0x0, 0x0, 0x0, 0x0}, v2_double = {0x0, 0x0}, v16_int8 = {0x0 <repeats 16 times>}, v8_int16 = {0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0},
v4_int32 = {0x0, 0x0, 0x0, 0x0}, v2_int64 = {0x0, 0x0}, uint128 = 0x0}
xmm5           {v4_float = {0x0, 0x0, 0x0, 0x0}, v2_double = {0x0, 0x0}, v16_int8 = {0x0 <repeats 16 times>}, v8_int16 = {0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0, 0x0},
v4_int32 = {0x0, 0x0, 0x0, 0x0}, v2_int64 = {0x0, 0x0}, uint128 = 0x0}
```

x/24x \$esp gives the content of the stack.

```
(gdb)
(gdb) x/24x $esp
0x7bdc: 0x00000000 0x00000000 0x00000000 0x00000000
0x7bec: 0x00000000 0x00000000 0x00000000 0x00000000
0x7bfc: 0x000007c4d 0x8ec031fa 0x8ec08ed8 0xa864e4d0
0x7c0c: 0xb0fa7502 0xe464e6d1 0x7502a864 0xe6dffb0fa
0x7c1c: 0x16010f60 0x200f7c78 0xc88366c0 0xc0220f01
0x7c2c: 0x087c31ea 0x10b86600 0x8ed88e00 0x66d08ec0
(gdb)
(gdb) |
```

Solution.

Explaining the non-zero values on the stack:

- 0x7bdc: This is the return address
- 0x7bfc: This is saved registers, instructions
- 0x7c0c: These have instructions
- 0x7c1c: These have instructions
- 0x7c2c: These have instructions

As seen above, The part of the stack printout that is actually the stack is from 0x7bdc to 0x7c2c (non-zero values). Since, it is the stack printout because it is not empty and contains data that has been pushed onto the stack during program execution. The other parts of the printout are mostly

empty, because they have not yet been used by the program.

Explaining the contents of the stack:

1. In bootasm.S, the stack pointer is initialized in the `_start` function. This initialization is done by the instruction `movl %esp, %ebp`. This instruction copies the value of the stack pointer `%esp` to the base pointer register, `%ebp`, which we are using to access function arguments and local variables.

2. For single step call to `bootmain`, we find its starting address through `objdump`.

```
gautam-ahuja@LAPTOP-FV7627LB:~/.../cs1217-assignment-3-julius-stabs-back$ objdump -d bootblock.o

bootblock.o:      file format elf32-i386

Disassembly of section .text:

00007c00 <start>:
 7c00:    fa                cli
 7c01:    31 c0              xor     %eax,%eax
 7c03:    8e d8              mov     %eax,%ds
 7c05:    8e c0              mov     %eax,%es
 7c07:    8e d0              mov     %eax,%ss

00007c09 <seta20.1>:
 7c09:    e4 64              in      $0x64,%al
 7c0b:    a8 02              test    $0x2,%al
 7c0d:    75 fa              jne     7c09 <seta20.1>
 7c0f:    b0 d1              mov     $0xd1,%al
 7c11:    e6 64              out     %al,$0x64
```

After single step through the call to `bootmain`, the stack will now have the return address and the arguments to call `bootmain`. The layout of the stack will look like this:

```
(gdb) br * 0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/24x $esp
0x6f00: 0xf000d009      0x00000000      0x00006f5e      0x00008148
0x6f10: 0x00008189      0x00000000      0x00000000      0x00008148
0x6f20: 0x00006f5e      0x00007cc5      0x00000000      0x00000000
0x6f30: 0x00006f5e      0x00000000      0x00006210      0x00007c00
0x6f40: 0x00000080      0x00000000      0x00000000      0x0000ee5bf
0x6f50: 0x000f1d3c      0x00000000      0x00007c00      0x00000000
(gdb) si
[ 0:7c01] => 0x7c01: xor     %eax,%eax
0x00007c01 in ?? ()
(gdb) si
[ 0:7c03] => 0x7c03: mov     %eax,%ds
0x00007c03 in ?? ()
(gdb) si
[ 0:7c05] => 0x7c05: mov     %eax,%es
0x00007c05 in ?? ()
(gdb) si
[ 0:7c07] => 0x7c07: mov     %eax,%ss
```

0x7c00: 0x8ec031fa	0x8ec08ed8	0xa864e4d0	0xb0fa7502
0x7c10: 0xe464e6d1	0x7502a864	0xe6dfb0fa	0x16010f60
(gdb)			
0x7c20: 0x200f7c78	0xc88366c0	0xc0220f01	0x087c31ea
0x7c30: 0x10b86600	0x8ed88e00	0x66d08ec0	0x8e0000b8
0x7c40: 0xbce88ee0	0x00007c00	0x0000f0e8	0x00b86600
0x7c50: 0xc289668a	0xb866ef66	0xef668ae0	0x9066feeb
0x7c60: 0x00000000	0x00000000	0x0000ffff	0x00cf9a00
0x7c70: 0x0000ffff	0x00cf9200	0x7c600017	0xf7ba0000
(gdb)			
0x7c80: 0xec000001	0x3cc0e083	0xc3f87540	0x57e58955
0x7c90: 0x0c5d8b53	0xfffffe5e8	0x0001b8ff	0xf2ba0000
0x7ca0: 0xee000001	0x0001f3ba	0xeed88900	0xe8c1d889
0x7cb0: 0x01f4ba08	0x89ee0000	0x10e8c1d8	0x0001f5ba
0x7cc0: 0xd889ee00	0x8318e8c1	0xf6bae0c8	0xee000001
0x7cd0: 0x000020b8	0x01f7ba00	0xe8ee0000	0xffffffff9e
0x7be0: 0x00000000	0x00000000	0x00000000	0x00000000
0x7bf0: 0x00000000	0x00000000	0x00000000	0x00000000
0x7c00: 0x8ec031fa	0x8ec08ed8	0xa864e4d0	0xb0fa7502
0x7c10: 0xe464e6d1	0x7502a864	0xe6dfb0fa	0x16010f60
(gdb)			
0x7c20: 0x200f7c78	0xc88366c0	0xc0220f01	0x087c31ea
0x7c30: 0x10b86600	0x8ed88e00	0x66d08ec0	0x8e0000b8
0x7c40: 0xbce88ee0	0x00007c00	0x0000f0e8	0x00b86600
0x7c50: 0xc289668a	0xb866ef66	0xef668ae0	0x9066feeb
0x7c60: 0x00000000	0x00000000	0x0000ffff	0x00cf9a00

3. The first assembly instructions of the function perform stack setup using the following instructions:

```

ASM bootblock.asm
229 // Read 'count' bytes at 'offset' from kernel into physical address 'pa'.
230 // Might copy more than asked.
231 void
232 readseg(uchar* pa, uint count, uint offset)
233 {
234     7cf4: 55                push    %ebp
235     7cf5: 89 e5            mov     %esp,%ebp
236     7cf7: 57              push    %edi
237     7cf8: 56              push    %esi
238     7cf9: 53              push    %ebx
239     7cfa: 83 ec 0c        sub     $0xc,%esp
240     7cfd: 8b 5d 08        mov     0x8(%ebp),%ebx
241     7d00: 8b 75 10        mov     0x10(%ebp),%esi
242     uchar* epa;
243

```

- push %ebp:** This instruction pushes the current value of the base pointer (ebp) onto the stack. The current frame pointer is saved with this step.
- mov %esp, %ebp:** This instruction moves the current value of the stack pointer (ESP) into the base pointer (EBP) register.
- sub \$0xc, %esp:** This instruction subtracts 8 from the stack pointer (esp) register. It save 8 bytes of space on the stack for local variables in the bootmain.

1. Exercise 1

Done by: Gautam Ahuja

Problem: To print out a line for each system call invocation with System Call name and return value.

Working:

As given in the question, we looked at `syscall()` function in `syscall.c`

```
void
syscall(void)
{
    int num;
    struct proc *curproc = myproc();

    num = curproc->tf->eax;
    if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
        curproc->tf->eax = syscalls[num]();
    } else {
        cprintf("%d %s: unknown sys call %d\n",
            curproc->pid, curproc->name, num);
        curproc->tf->eax = -1;
    }
}
```

Here we see that a function `myproc()` is storing its results in another structure `tf` (trapframe) within the `proc` structure.

Looking at the definition, we can see that `proc` structure captures the current state of ongoing process.

```
// Per-process state
github-classroom[bot], 5 days ago | 1 author (github-classroom[bot])
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // switch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;       // Current directory
    char name[16];          // Process name (debugging)
};
```

Within this, the `tf` (trapframe) structure holds the information of all the registers for a process (defined in `x86.h`).

```

struct trapframe {
    // registers as pushed by pusha
    uint edi;
    uint esi;
    uint ebp;
    uint oesp;    // useless & ignored
    uint ebx;
    uint edx;
    uint ecx;
    uint eax;

    // rest of trap frame
    ushort gs;
    ushort padding1;
    ushort fs;
    ushort padding2;
    ushort es;
    ushort padding3;
    ushort ds;
    ushort padding4;
    uint trapno;

    // below here defined by x86 hardware
    uint err;
    uint eip;
    ushort cs;
    ushort padding5;
    uint eflags;

    // below here only when crossing rings, such as from user to kernel
    uint esp;
    ushort ss;
    ushort padding6;
};

```

Here the `eax` (known as "accumulator") register holds the the return value of `myproc()`.

After that it checks if the value is valid (> 0) and is within the range of defined number of syscalls and a that a syscall of `num` exists.

If result is true, it calls the particular function numbered (which are numbered in `syscall.h`) in a functional structure of `syscalls` defined above. The resulting return value is stored in `eax` register again of above defined `cpuproc` structure.

```

C syscall.h
1 // System call numbers
2 #define SYS_fork 1
3 #define SYS_exit 2
4 #define SYS_wait 3
5 #define SYS_pipe 4
6 #define SYS_read 5
7 #define SYS_kill 6
8 #define SYS_exec 7
9 #define SYS_fstat 8
10 #define SYS_chdir 9
11 #define SYS_dup 10
12 #define SYS_getpid 11
13 #define SYS_sbrk 12
14 #define SYS_sleep 13
15 #define SYS_uptime 14
16 #define SYS_open 15
17 #define SYS_write 16
18 #define SYS_mknod 17
19 #define SYS_unlink 18
20 #define SYS_link 19
21 #define SYS_mkdir 20
22 #define SYS_close 21

```

Solution:

We defined a static array of pointers (string literals) which contains the syscall names corresponding to their defined number in `syscall.h` as follows:

Then we just run a switch-case inside the `if` statement of `syscall()` which prints out the syscall name and its corresponding number currently in `eax` register as per current `num`.

The result of above is:

```

$XDP (https://lpxe.org/) 00:03:0 CA00 PC12.10 PnP PMM+IF8B9A40+1FECB9A0 CA00
Bootlog (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>

This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word".
* target remote localhost:26980
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
The target architecture is set to "i8086".
[ro000ffff] <error: lyp $ox3c30,$oxf000eb5b
&xvmmrirs in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configu
ration
of GDB. Attempting to continue with the default i8086 settings.

(gdb) br = 0x0010000c
Breakpoint 1 at 0x10000c
(gdb) c
Continuing.
The target architecture is set to "i8086".
=> 0x10000c: mov %cr4,%eax

Thread 1 hit breakpoint 1, 0x0010000c in ?? ()
(gdb) c
Continuing.
```


2. Exercise 2

Done by: Gautam Ahuja

Problem: To add a new system call to xv6 for printing date.

Working:

As given in the question we run `grep -n uptime *.c` to clone all of the pieces of code related to uptime.

```
gautam-ahuja@LAPTOP-FV7627LB:~/../cs1217-assignment-3-julius-stabs-back$ grep -n uptime *.c
syscall.c:106:extern int sys_uptime(void);
syscall.c:124:[SYS_uptime] sys_uptime,
syscall.c:149: [SYS_uptime] "uptime",
syscall.c:212:     case SYS_uptime:
syscall.c:301://         cprintf("uptime -> %d\n", curproc->tf->eax);
syscall.h:15:#define SYS_uptime 14
sysproc.c:83:sys_uptime(void)
user.h:25:int uptime(void);
usys.S:31:SYSCALL(uptime)
gautam-ahuja@LAPTOP-FV7627LB:~/../cs1217-assignment-3-julius-stabs-back$ |
```

From here we get an idea about the files we need to modify: `syscall.c`, `syscall.h`, `user.h`, `usys.S`, `sysproc.c`, `Makefile`.

Next we create a `date.c` file with the code:

```
#include "types.h"
#include "user.h"
#include "date.h"
int
main(int argc, char *argv[])
{
    struct rtcdate r;
    if (date(&r)) {
        // 2 is file descriptor for stderr
        printf(2, "date failed\n");
        exit();
    }
    // 1 is file descriptor for stdout
    printf(1, "Year: %d\nMonth: %d\nDay: %d\nHour : Minute : Seconds :: %d:%d:%d\n",
           r.year, r.month, r.day, r.hour, r.minute, r.second);
    exit();
}
```

In above, the first `printf()` statement inside `if` condition writes to file descriptor 2, which is `stderr`.

The second `printf()` statement writes to output, through file descriptor 1, which is `stdout`.

Next, we defined a new syscall in `syscall.h` by `#define SYS_date 22` and give it a number 22.

```
C syscall.h
1 // System call numbers
2 #define SYS_fork 1
3 #define SYS_exit 2
4 #define SYS_wait 3
5 #define SYS_pipe 4
6 #define SYS_read 5
7 #define SYS_kill 6
8 #define SYS_exec 7
9 #define SYS_fstat 8
10 #define SYS_chdir 9
11 #define SYS_dup 10
12 #define SYS_getpid 11
13 #define SYS_sbrk 12
14 #define SYS_sleep 13
15 #define SYS_uptime 14
16 #define SYS_open 15
17 #define SYS_write 16
18 #define SYS_mknod 17
19 #define SYS_unlink 18
20 #define SYS_link 19
21 #define SYS_mkdir 20
22 #define SYS_close 21
23 #define SYS_date 22
```

We also define the function type in `user.h`. Since the function which takes in argument a pointer to `rtcddate` structure. It is defined as: `int date(struct rtcddate*);`.

```
char *sbrk(int);
int sleep(int);
int uptime(void);
// adding a new system call of date, it takes a pointer to a rtcddate struct as an argument
int date(struct rtcddate*);
// ...
```

Next we add `_date\` to `UPROGS` which is a list of user-level programs in the Makefile.

```
UPROGS=\
_cat\
_echo\
_forktest\
_grep\
_init\
_kill\
_ln\
_ls\
_mkdir\
_rm\
_sh\
_stressfs\
_usertests\
_wc\
_zombie\
_date\
```

We also define a new syscall to the assembly file for user calls (`usys.S`) as:

```

asm usys.S
19 SYSCALL(exec)
20 SYSCALL(open)
21 SYSCALL(mknod)
22 SYSCALL(unlink)
23 SYSCALL(fstat)
24 SYSCALL(link)
25 SYSCALL(mkdir)
26 SYSCALL(chdir)
27 SYSCALL(dup)
28 SYSCALL(getpid)
29 SYSCALL(sbrk)
30 SYSCALL(sleep)
31 SYSCALL(uptime)
32 SYSCALL(date)
33 |

```

Lastly we defined the implementation of date syscall in `sysproc.c` which includes system calls that are implemented in relation to management of processes. The `trapframe` discusses in exercise 1 always looks for a definition of a syscall function in `sysproc.c` when it encounters a syscall.

```

C sysproc.c
92
93 // Date edits
94 int
95 sys_date(void)
96 {
97     struct rtcdate *r;
98     if(argptr(0, (void*)&r, sizeof(*r)) < 0)
99         return -1;
100     cmostime(r);
101     return 0;
102 }

```

The `if` statement checks if the incoming pointer to syscall is valid or not. The `argptr` takes input (through file descriptor 0, which is `stdin`) and checks if it lies within memory space. If not, it exits the syscall with a return code of `-1`. The definition of `argptr` is:

```

int argptr(int, char **, int)
Fetch the nth word-sized system call argument as a pointer
to a block of memory of size bytes. Check that the pointer
lies within the process address space.

```

At last we add `sys_date` to the `syscall.c` file to the functional array declaration.

```

C syscalls.c
108 // date calls
109
110 static int (*syscalls[])(void) = {
111     [SYS_fork]    sys_fork,
112     [SYS_exit]    sys_exit,
113     [SYS_wait]    sys_wait,
114     [SYS_pipe]    sys_pipe,
115     [SYS_read]    sys_read,
116     [SYS_kill]    sys_kill,
117     [SYS_exec]    sys_exec,
118     [SYS_fstat]   sys_fstat,
119     [SYS_chdir]   sys_chdir,
120     [SYS_dup]     sys_dup,
121     [SYS_getpid]  sys_getpid,
122     [SYS_sbrk]    sys_sbrk,
123     [SYS_sleep]   sys_sleep,
124     [SYS_uptime]  sys_uptime,
125     [SYS_open]    sys_open,
126     [SYS_write]   sys_write,
127     [SYS_mknod]   sys_mknod,
128     [SYS_unlink]  sys_unlink,
129     [SYS_link]    sys_link,
130     [SYS_mkdir]   sys_mkdir,
131     [SYS_close]   sys_close,
132     [SYS_date]    sys_date,
133 };

```

The output of the date syscall is:

```

SeaBIOS (version 1.15.0-1)

iPXE (https://ipxe.org) 00:03:0 CA00 PCI2.10 PnP PMM+1FF8B4A0+1FECB4A0 CA00

Booting from Hard Disk..xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap8
init: starting sh
$ date
Year: 2023
Month: 2
Day: 20
Hour : Minute : Seconds :: 16:3:29
$

gautam-ahuja@LAPTOP-FV76:~/.../cs1217-assignment-3-julius-stabs-back$ gd
GNU gdb (Ubuntu 12.1-0ubuntu1-22.04) 12.1
Copyright (C) 2022 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<https://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.

For help, type "help".
Type "apropos word" to search for commands related to "word".
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
The target architecture is set to "i8086".
[f000:ffff] 0xfffff0: jmp $0x3630,$0xf000e05b
0x0000ffff in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configu
ration
of GDB. Attempting to continue with the default i8086 settings.

(gdb) br * 0x0010000c
Breakpoint 1 at 0x10000c
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x10000c: mov %cr4,%eax

Thread 1 hit Breakpoint 1, 0x0010000c in ?? ()
(gdb) c
Continuing.

```