

CS1217 - Spring 2023 - Lab 2

Gautam Ahuja, Nistha Singh

01. Exercise 01

The various states that an `xv6` process can be found in the `procstate` enum data type in the `proc.h` file. These are:

- **Unused:** This is the initial state of the process. It is not being used, executed or performing any tasks.
- **Embryo:** This is the state of the process when it is in formation. For example, when a `fork()` is called, the state in which the process is being created, its variables are being set, address state is being copied, etc is the **Embryo** state.
- **Sleeping:** It is the waiting/blocking/sleeping state of a process. It occurs when the process is waiting on some other event to happen, or a process calls `sleep()`, etc.
- **Runnable:** This is equivalent to the ready state. The process in this state can start executing as soon as it gets hold of the CPU.
- **Running:** It is the process which is currently executing instructions on the CPU.
- **Zombie:** This state is when a process is terminated (is finished executing). This can also be a process which has no parent process.

There are a total of **6 states** that an `xv6` process can be in. These states are used to manage the process by the operating system `xv6`.

```
C:proc.h > ...
35 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
36
```

The `state` variable of `enum procstate` type keeps the track of the current state of the process inside the `proc` data structure.

```
C:proc.h > proc > state
37 // Per-process state
38 struct proc {
39     uint sz; // Size of process memory (bytes)
40     pde_t* pgdir; // Page table
41     char *kstack; // Bottom of kernel stack for this process
42     enum procstate state; // Process state
43     int pid; // Process ID
44     struct proc *parent; // Parent process
45     struct trapframe *tf; // Trap frame for current syscall
46     struct context *context; // switch() here to run process
47     void *chan; // If non-zero, sleeping on chan
48     int killed; // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd; // Current directory
51     char name[16]; // Process name (debugging)
52 };
53
```

02. Exercise 02

The scheduler code is:

```

C procc > scheduler(void)
322 void
323 scheduler(void)
324 {
325     struct proc *p;
326     struct cpu *c = mycpu();
327     c->proc = 0;
328
329     for(;;){
330         // Enable interrupts on this processor. | github-classroom[bot], last week * Initial commit
331         sti(); // STI - Set Interrupt Flag
332
333         // Loop over process table looking for process to run.
334         acquire(&ptable.lock); // the &ptable.proc[NPROC] is the address of the last element of the ptable.proc array.
335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
336             if(p->state != RUNNABLE)
337                 continue;
338
339             // Switch to chosen process. It is the process's job
340             // to release ptable.lock and then reacquire it
341             // before jumping back to us.
342             c->proc = p;
343             switchvm(p); // the switchvm(p) will do the switching of the kernel virtual memory to the process's virtual memory.
344             p->state = RUNNING; // Now the process is running, so we can switch to it.
345             // The switch(&(c->scheduler), p->context) is the context switch. It will save the current context and load the new context.
346             swtch(&(c->scheduler), p->context);
347             switchvm(); // The switchvm function will switch the process's virtual memory to the kernel virtual memory.
348
349             // Process is done running for now.
350             // It should have changed its p->state before coming back.
351             c->proc = 0;
352         }
353         release(&ptable.lock);
354     }
355 }

```

Line 329 has a `for` loop that looks like an infinite loop. This is done to make sure the `scheduler` always runs. As we know, the scheduler is an essential algorithm that makes sure the CPU is used at all the time.

Since the entire logic of scheduler is inside the loop, it makes sure that the kernel can keep track of all `RUNNABLE` processes.

Without the infinite loop, the scheduler would run only once and kernel will not be able to put any more processes on CPU after one exits.

The implemented scheduler policy is a **Round Robin Scheduling Algorithm**.

As discussed in class, a round-robin policy maintains a list of processes and iterates over to execute every runnable process.

In above algorithm, line **334** acquired the lock of the page-table (`ptable`) which holds the list of all the processes (`ptable.proc`). The maximum number of processes can only be 64 as defined by `NPROC` inside `param.h`.

After this, the condition at line **336** checks if the process is in runnable state, if not, increment the pointer `p` to point to next process on the list. After identifying the runnable process, the scheduler puts in on the CPU (`c->proc = p;`) on line **342** and changes the state from `RUNNABLE` to `RUNNING` on line **344**.

The kernel also switches to the processes virtual address space via `switchvm(p);` on line **343** which loads the process's page table and then do a context switch via `swtch(&(c->scheduler), p->context);` on line **346**. At this point the new user process is running on the CPU until an event occurs (`exit()`, `yield()`, `wait()`, `syscall`, or timer interrupt). The definition of `swtch()` can

be found in the `switch.S` file.

```

ASM switch.S
github-classroom[bot], last week | 1 author (github-classroom[bot])
1  # Context switch
2  #
3  # void switch(struct context **old, struct context *new);
4  #
5  # Save the current registers on the stack, creating
6  # a struct context, and save its address in *old.
7  # Switch stacks to new and pop previously-saved registers.
8

```

At this point (after completion) the scheduler switched the memory back from process' view to kernel-only global page table view (line **347** - `switchkvm()`;). The CPU's current running process is set to NULL and the loop runs again to put the next `RUNNABLE` process onto the CPU.

In case the process is not completed in time, the time interrupt will occur which will be identified by `sti()`;. After addressing the IRS the context will switch to kernel mode where the scheduler will run again from line **347**, the kernel address space, and continue to choose the next `RUNNABLE` process.

In such way the **Round Robin Scheduling policy** is implemented.

the definition of `switchuvm()` and `switchkvm()` is as follows:

```

void switchuvm(struct proc *)
Switch TSS and h/w page table to correspond to process p.
switchuvm(p);           // the switchuvm(p) will

```

```

void switchkvm(void)
Switch h/w page table register to the kernel-only page table,
for when no process is running.
switchkvm();           // The switchkvm function

```

03. Exercise 03

The definition of `fork()` is:

```

C proc > fork(void)
176
177 // Create a new process copying p as the parent.
178 // Sets up stack to return as if from system call.
179 // Caller must set state of returned proc to RUNNABLE.
180 int
181 fork(void)
182 {
183     int i, pid;
184     struct proc *np;
185     struct proc *curproc = myproc();
186     // Allocate process.
187     if((np = allocproc()) == 0){
188         return -1;
189     }
190     // Copy process state from proc.
191     if((np->pgdir = copyuvm(curproc->pgdir, curproc->sz)) == 0){
192         kfree(np->kstack);
193         np->kstack = 0;
194         np->state = UNUSED;
195         return -1;
196     }
197     np->sz = curproc->sz;
198     np->parent = curproc;
199     *np->tf = *curproc->tf;
200     // Clear %eax so that fork returns 0 in the child.
201     np->tf->eax = 0;
202     for(i = 0; i < NOFILE; i++)
203         if(curproc->ofile[i])
204             np->ofile[i] = filedup(curproc->ofile[i]);
205     np->cwd = idup(curproc->cwd);
206     safestrcpy(np->name, curproc->name, sizeof(curproc->name));
207     pid = np->pid;
208     acquire(&table.lock);
209     np->state = RUNNABLE;
210     release(&table.lock);
211     return pid;
212 }
213

```

The `fork()` is executed in an order of:

1. It first checks if there is any slots available in the `ptable` through the `allocproc()` function; which allocates and `=` returns the pointer to new `proc` of new function, if available, else returns 0; in which case the `fork()` return -1.

```

static struct proc *allocproc(void)
{
    PAGEBREAK: 32
    Look in the process table for an UNUSED proc.
    If found, change state to EMBRYO and initialize
    state required to run in the kernel.
    Otherwise return 0.
    allocproc() == 0){

```

2. It also copies over the parent process's memory space to the child process's memory space with `copyuvm(curproc->pgdir, curproc->sz)`. If the copying over was unsuccessful (in which case `copyuvm()` returns 0), the `fork()` return -1 and the state of the process is set as `UNUSED`.

```

pde_t *copyuvm(pde_t *, uint)
    Given a parent process's page table, create a copy
    of it for a child.
    copyuvm(curproc->pgdir, curproc->sz) == 0)

```

When the process of copying of address space is executing, the child is in the state of `EMBRYO`

3. In normal scenario where the address space is copied successfully, the stack pointer for child is set, `pid` of parent is recorded and the process is put into a `RUNNABLE` state as the child is now ready to be executed. The `fork()` returns the `pid` of the child in parent process.

04. Exercise 04

As discussed in Question 1, a ZOMBIE process is the one which has no parent or in other words it is abandoned.

As discussed in class, the abandoned are reassigned to the grandfather (if available) or the `init` process itself. Doing a search for all occurrences of keyword ZOMBIE in `proc.c` we find out that the process `exit()` is used to reassign the abandoned processes to the `init` process.

```

C proc.c > ...
227 void
228 exit(void)
229 {
230     struct proc *curproc = myproc();
231     struct proc *p;
232     int fd;
233     if(curproc == initproc)
234         panic("init exiting");
235     // Close all open files.
236     for(fd = 0; fd < NOFILE; fd++){
237         if(curproc->ofile[fd]){
238             fileclose(curproc->ofile[fd]);
239             curproc->ofile[fd] = 0;
240         }
241     }
242     begin_op();
243     input(curproc->cwd);
244     end_op();
245     curproc->cwd = 0;
246     acquire(&ptable.lock);
247     // Parent might be sleeping in wait().
248     wakeup1(curproc->parent);
249     // Pass abandoned children to init.
250     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
251         if(p->parent == curproc){
252             p->parent = initproc;
253             if(p->state == ZOMBIE)
254                 wakeup1(initproc);
255         }
256     }
257     // Jump into the scheduler, never to return.
258     curproc->state = ZOMBIE;
259     sched();
260     panic("zombie exit");
261 }
262

```

Here we can see in the section of lines **249-256**, if a process is abandoned (ZOMBIE), it is being reassigned to the `init` function itself.

I. Part 01: MLFQ Scheduler

Final File: *proc_mlfq_v02.c*

As done in assignment 3, we know to add a new system call we need to modify the following files: `syscall.c`, `syscall.h`, `user.h`, `usys.S`, `sysproc.c`, `Makefile`.

We will also need to edit the following files: `proc.c`, `proc.h`, `param.h` as per the question.

To add a new system call, we first define it in `syscall.h` as follows

```
C syscall.h > ...
22 #define SYS_close 21
23 // Two new system calls - start
24 #define SYS_setpriority 22
25 #define SYS_getpriority 23
26 // Two new system calls - end
```

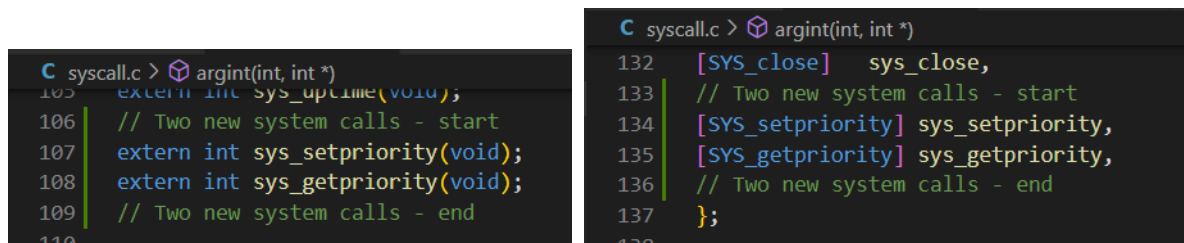
We also add `MAXPRIORITY` and `DEFAULT_BUDGET` in files `proc.h` as per the question:

```
C proc.h > DEFAULT_BUDGET
2 // Defining the maxpriority and max budget
3 #define MAXPRIORITY 2
4 #define DEFAULT_BUDGET 5
5 #define TICKS_TO_PROMOTE 20
```

We also add a new parameter of `priority` and `budget` in the `proc` structure in `proc.h` to hold the priority of current process.

```
C proc.h > ...
62 struct inode *cwd; // Current directory
63 char name[16]; // Process name (debugging)
64 // Lab 2
65 int priority; // Process priority
66 // int budget; // Process budget
67 };
```

Then in `syscall.c` we add a reference to each system call function and we also add it to the system call table which maps the functions to system call numbers defined in `syscall.h`



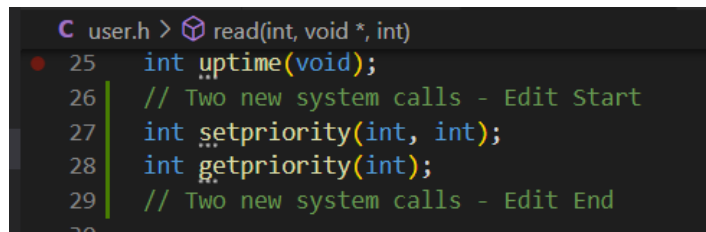
```

C syscall.c > argint(int, int *)
105 extern int sys_uptime(void);
106 // Two new system calls - start
107 extern int sys_setpriority(void);
108 extern int sys_getpriority(void);
109 // Two new system calls - end
110

C syscall.c > argint(int, int *)
132 [SYS_close] sys_close,
133 // Two new system calls - start
134 [SYS_setpriority] sys_setpriority,
135 [SYS_getpriority] sys_getpriority,
136 // Two new system calls - end
137 };

```

The declaration of functions (to which the syscall will call) is added in the file `user.h`

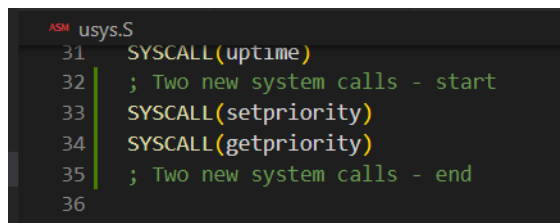


```

C user.h > read(int, void *, int)
25 int uptime(void);
26 // Two new system calls - Edit Start
27 int setpriority(int, int);
28 int getpriority(int);
29 // Two new system calls - Edit End
30

```

The declaration of these functions are also added in the assembly for user system calls in the file `usys.S`

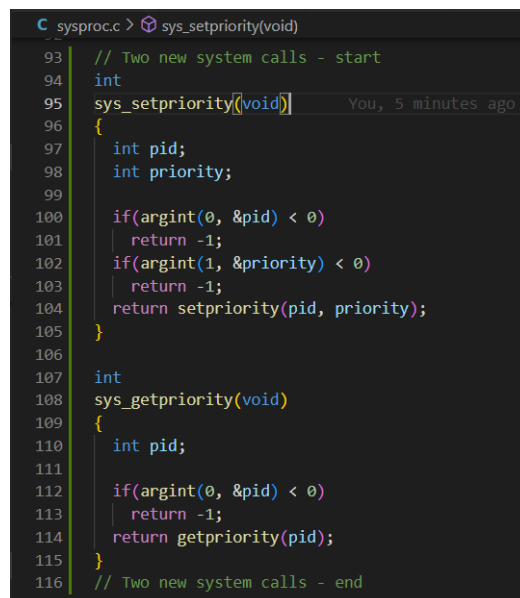


```

ASM usys.S
31 SYSCALL(uptime)
32 ; Two new system calls - start
33 SYSCALL(setpriority)
34 SYSCALL(getpriority)
35 ; Two new system calls - end
36

```

The definition of system call functions - `sys_getpriority()` and `sys_setpriority()` (which will call to the actual function - `getpriority()` and `setpriority()`) is implemented in `sysproc.c`



```

C sysproc.c > sys_setpriority(void)
93 // Two new system calls - start
94 int
95 sys_setpriority([void]) You, 5 minutes ago
96 {
97     int pid;
98     int priority;
99
100     if(argint(0, &pid) < 0)
101         return -1;
102     if(argint(1, &priority) < 0)
103         return -1;
104     return setpriority(pid, priority);
105 }
106
107 int
108 sys_getpriority(void)
109 {
110     int pid;
111
112     if(argint(0, &pid) < 0)
113         return -1;
114     return getpriority(pid);
115 }
116 // Two new system calls - end

```

Finally the functions `getpriority()` and `setpriority()` are defined in `proc.c` as follows:

The `setpriority()` takes the `pid` of the process and sets the `priority` which is passed as an argument. It first checks if the `priority` is valid and if the process with given `pid` exists, it assigns the `priority` and sets the budget to `DEFAULT_BUDGET`. In other cases it returns an error code.

Similarly the `getpriority()` takes the `pid` as an argument. If the process `pid` is valid and the process is not in `UNUSED` state, it returns the `pid`. Else it returns error code in other cases.

```
C proc.c > setpriority(int, int)
536 // New system calls
537 // int setpriority(int pid, int priority)
538 int
539 setpriority(int pid, int priority)
540 {
541     struct proc *p;
542     int returnCode = -1; // returnCode 0 if success, -1 if fail
543
544     // check if priority is valid
545     if(priority < 0 || priority > MAXPRIORITY){
546         cprintf("Invalid priority value detected");
547         return returnCode;
548     }
549
550     acquire(&ptable.lock);
551     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
552         if(p->pid == pid){
553             p->priority = priority;
554             p->budget = DEFAULT_BUDGET;
555             returnCode = 0;
556             break;
557         }
558     }
559     release(&ptable.lock);
560
561     // check if pid is valid
562     if(returnCode == -1){
563         cprintf("Invalid pid value detected");
564     }
565
566     return returnCode;
567 }
```

```
C proc.c > getpriority(int)
569 // int getpriority(int pid)
570 int
571 getpriority(int pid)
572 {
573     struct proc *p;
574     int priority = -1;
575
576     acquire(&ptable.lock);
577     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
578         if(p->pid == pid && p->state != UNUSED){
579             priority = p->priority;
580             break;
581         }
582         if(p->pid == pid && p->state == UNUSED){
583             cprintf("Process with pid %d is in UNUSED state", pid);
584             return priority;
585         }
586     }
587     release(&ptable.lock);
588
589     // check if pid is valid
590     if(priority == -1){
591         cprintf("Invalid pid value detected");
592     }
593
594     // End of new system calls
```

At last we edit the `defs.h` to include all the changes and add the function definitions.

```
C defs.h
123 // Two new system calls - start
124 int      setpriority(int, int);
125 int      getpriority(int);
126 // Two new system calls - end
```

As per the question, we add a new field to `proc` structure which holds the original `priority` in case the process goes to any other state from `RUNNABLE`. It will help us determine where to put the process once it is `RUNNABLE` again.


```

C proc.h > proc
64 // Lab 2
65 int priority;           // Process priority
66 int budget;            // Process budget
67 int original_priority;  // Process original priority, keep track of if process state changes
68 };

```

We also edit the `fork()` in `proc.c` as upon allocation, each process will have the same initial (default) priority value, the highest priority. We set the process priority to `MAXPRRORITY`.

```

C proc.c > fork(void)
217 np->state = RUNNABLE;
218 // set the priority of the child process to highest priority
219 np->priority = MAXPRRORITY;
220

```

Periodic Priority Adjustment:

1. We define a new field of `PromoteAtTime` in the `ptable` structure. Which will store the ticks value at which promotion will occur.

```

C procc > ptable
14 struct {
15     struct spinlock lock;
16     struct proc proc[NPROC];
17     // timer variables
18     uint PromoteAtTime;
19 } ptable;

```

2. We set the page table value of `PromoteAtTime` in `userinit()` as follows:

```

C procc > userinit(void)
156 acquire(&ptable.lock);
157
158 p->state = RUNNABLE;
159 // set promotion time
160 ptable.PromoteAtTime = ticks + TICKS_TO_PROMOTE;
161
162 release(&ptable.lock);

```

3. We define `TICKS_TO_PROMOTE` as:

```

C proc.h > TICKS_TO_PROMOTE
4 #define TICKS_TO_PROMOTE 20

```

4. Checking priorities each time `scheduler()` runs:

Note: We did not use the `getpriority()` in the implementation of the scheduler. They themselves being a system call (which also calls `getpid()` system call) adds an overhead to the scheduler (and kernel in general) and thus slowing it down and sometimes not working at all.

We define a new 2D array of `struct proc`. Each rows is the priority level (3rd being the highest) and each row can contain up to 64 (`NPROC`) processes each. The `mlfq_queue_size` keeps track of the amount of non-ZOMBIE process on each levels

```

C proc_mlfq_v02.c > wakeup(void *)
10 // Define three levels of priority for MLFQ Scheduler
    You, 16 hours ago
11 struct proc *mlfq_queue[MAXPRIORITY+1][2*NPROC];
12 // top queue has highest priority
13 int mlfq_queue_size[MAXPRIORITY+1] = {-1,-1,-1};
14

```

To adjust the new process according to new priority we do as follows:

```

C proc_mlfq_v02.c > scheduler(void)
393 /**** Periodic Priority Adjustment: Part 3 ****/
394 if(ticks >= ptable.PromoteAtTime){
395     //cprintf("Promotion Time");
396     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
397         if(p->state != ZOMBIE && p->priority < MAXPRIORITY){
398             // delete the process from the queue at that priority level
399             for(int i = 0; i <= mlfq_queue_size[p->priority]; i++){
400                 if(mlfq_queue[p->priority][i] == p){
401                     for(int j = i; j < mlfq_queue_size[p->priority]; j++){
402                         mlfq_queue[p->priority][j] = mlfq_queue[p->priority][j+1];
403                     }
404                     mlfq_queue_size[p->priority]--;
405                     break;
406                 }
407             }
408             // p->priority++;
409             setpriority(p->pid, p->priority + 1);
410             p->budget = DEFAULT_BUDGET;
411             // // send it to the end of the next priority queue
412             mlfq_queue_size[p->priority]++;
413             mlfq_queue[p->priority][mlfq_queue_size[p->priority]] = p;
414         }
415     }
416     ptable.PromoteAtTime = ticks + TICKS_TO_PROMOTE;
417 }
    You, yesterday • MLFQ v1.0

```

We check if the `ticks` have been exceeded or not. Then going through all the processes, we skip the `ZOMBIE` processes. For every other process, we follow the logic of:

† We assume that max processes in top priority can be 64. Any more than that will not be possible.

- Pop the process out of the current priority level and shift all processes in that level left by one to fill the gap.
- Increase priority of this popped out process.
- Insert this new process at the end of new priority queue.

At the end, update the `PromoteAtTime` to `ticks + TICKS_TO_PROMOTE`.

We also edit the `fork()` so that every new child process created gets the `MAXPRIORITY` and `DEFAULT_BUDGET` allocated to it.

```

C proc.c > fork(void)
231 // set the priority of the child process to highest priority
232 np->priority = MAXPRIORITY;
233 np->original_priority = MAXPRIORITY;
234 // set the budget of the child process to default budget
235 np->budget = DEFAULT_BUDGET;

```

MLFQ Scheduler:

As discussed, whenever a new process is created, it is assigned a budget of `DEFAULT_BUDGET` and a priority of `MAXPRIORITY`.

```

C proc_mlfq_v02.c > fork(void)
239 np->state = RUNNABLE;
240 // set the priority of the child process to highest priority
241 np->priority = MAXPRIORITY;
242 np->original_priority = MAXPRIORITY;
243 // // set the budget of the child process to default budget
244 np->budget = DEFAULT_BUDGET;
245 // // New process created is inserted at the end (tail) of the queue
246 mlfq_queue_size[MAXPRIORITY]++;
247 mlfq_queue[MAXPRIORITY][mlfq_queue_size[MAXPRIORITY]] = np;
248 // cprintf("MAXPRIORITY queue size: %d\n", mlfq_queue_size[MAXPRIORITY]);

```

We also edit the `wakeup1()` function. Whenever a process is awoken, it is checked if it belongs to the priority queue as per the `p->priority` of the state. If not, we send it to the tail of the priority queue it belongs to.

```
C proc_mlfq_v02.c > wakeup1(void *)
606 static void
607 wakeup1(void *chan)
608 {
609     struct proc *p;
610
611     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
612         if(p->state == SLEEPING && p->chan == chan){
613             p->state = RUNNABLE;
614             // find if already in priority queue
615             int check = 0;
616             for(int i = 0; i < mlfq_queue_size[p->priority]; i++){
617                 if(mlfq_queue[p->priority][i] == p){
618                     check = 1;
619                     break;
620                 }
621             }
622             if(check == 0){
623                 mlfq_queue_size[p->priority]++;
624                 mlfq_queue[p->priority][mlfq_queue_size[p->priority]] = p;
625             }
626 }
```

For the actual implementation of the MLFQ Scheduling algorithm, we have our 2D array with rows as priority levels and columns as the process entry in those priority queues.

We simply iterate through each level (startig with `MAXPRIORITY`) untill its all the processes in that level are finished (size of that queue becomes less than zero).

For every `RUNNABLE` process, we put it on the `cpu` and note down the time of its execution.

```
C proc_mlfq_v02.c > scheduler(void)
423 // ***** MLFQ Scheduler *****
424 for(int i = MAXPRIORITY; i >= 0; i--){
425     if(mlfq_queue_size[i] < 0){
426         continue;
427     }
428     while(mlfq_queue_size[i] >= 0){
429         int check = 0;
430         for(int k=0; k <= mlfq_queue_size[i]; k++){
431             if(mlfq_queue[i][k]->state != RUNNABLE && k==mlfq_queue_size[i]){
432                 break;
433             }
434             if(mlfq_queue[i][k]->state != RUNNABLE){
435                 continue;
436             }
437             p = mlfq_queue[i][k];
438             // remove the process from the queue
439             for(int temp = k; temp < mlfq_queue_size[i]; temp++){
440                 mlfq_queue[i][temp] = mlfq_queue[i][temp+1];
441             }
442             mlfq_queue_size[i]--;
443             check = 1;
444             // Note the start time of the process
445             int start_time = ticks; //uptime();
446             c->proc = p;
447             switchvm(p);
448             p->state = RUNNING;
449             swtch(&(c->scheduler), p->context);
450             switchvm();
451             c->proc = 0;
452             int end_time = ticks; //uptime();
453             // Update the budget of the process
454             p->budget = p->budget - (end_time - start_time);
```

We update the budget of each program as: `p->budget = p->budget - (end.time - start.time)`. Now as per given condition if `budget < 0`, we demote the process to lower priority level and insert it at the tail of the queue. Else we move it to the tail of current priority level.

```

C proc_mlfq_v02.c > scheduler(void)
454     p->budget = p->budget - (end_time - start_time);
455     if(p->budget<=0){
456         //process already removed from current queue
457         // demote priority and place it at the tail of the new queue
458         if(p->priority > 0){
459             setpriority(p->pid, p->priority - 1);
460             p->original_priority = p->priority;
461             mlfq_queue_size[p->priority]++;
462             mlfq_queue[p->priority][mlfq_queue_size[p->priority]] = p;
463         }
464         else{
465             // if priority is 0 then send to the end of the queue
466             mlfq_queue_size[p->priority]++;
467             mlfq_queue[p->priority][mlfq_queue_size[p->priority]] = p;
468         }
469     }
470     else{
471         //if process id still RUNNABLE then put it at the end of the queue
472         mlfq_queue_size[p->priority]++;
473         mlfq_queue[p->priority][mlfq_queue_size[p->priority]] = p;
474     }
475     }
476     if(check==0){
477         break;
478     }
479 }
480
481 release(&ptable.lock);
482 }
483 }
484

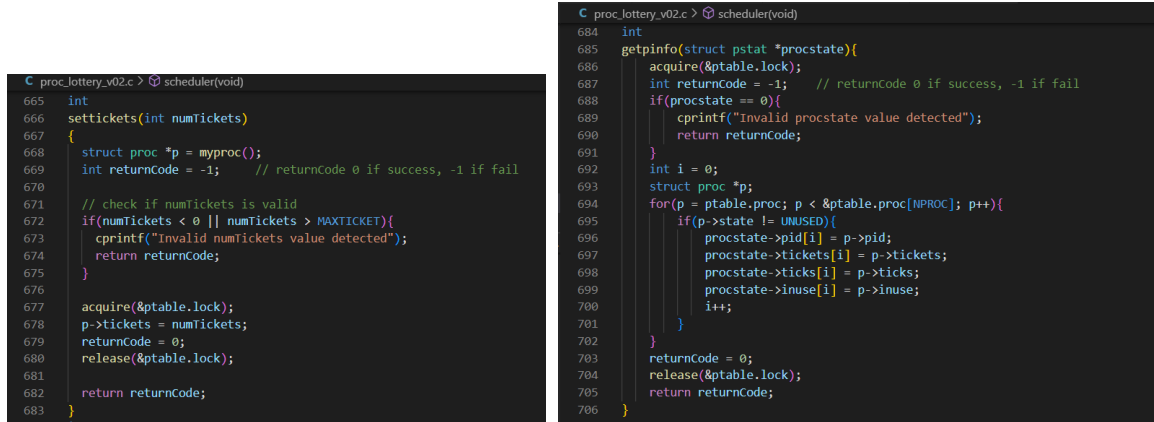
```

Along with periodic priority updating, the MLFQ Scheduling policy is implemented in xv6 operating system.

II. Part 02: Lottery Scheduler

Final File: *proc_lottery_v02.c*

As per the question, we start off by implementing two new system calls by editing a bunch of files (same as in previous parts): `setticket()` and `getpinfo()`. The definition of these functions are as follows:



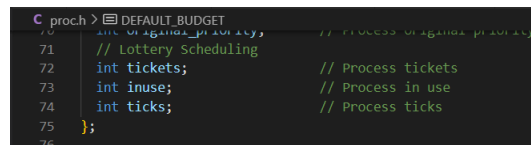
```

C proc_lottery_v02.c > scheduler(void)
665 int
666 settickets(int numtickets)
667 {
668     struct proc *p = myproc();
669     int returnCode = -1; // returnCode 0 if success, -1 if fail
670
671     // check if numtickets is valid
672     if(numtickets < 0 || numtickets > MAXTICKET){
673         cprintf("Invalid numtickets value detected");
674         return returnCode;
675     }
676
677     acquire(&ptable.lock);
678     p->tickets = numtickets;
679     returnCode = 0;
680     release(&ptable.lock);
681
682     return returnCode;
683 }

C proc_lottery_v02.c > scheduler(void)
684 int
685 getpinfo(struct pstat *procstate){
686     acquire(&ptable.lock);
687     int returnCode = -1; // returnCode 0 if success, -1 if fail
688     if(procstate == 0){
689         cprintf("Invalid procstate value detected");
690         return returnCode;
691     }
692     int i = 0;
693     struct proc *p;
694     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
695         if(p->state != UNUSED){
696             procstate->pid[i] = p->pid;
697             procstate->tickets[i] = p->tickets;
698             procstate->ticks[i] = p->ticks;
699             procstate->inuse[i] = p->inuse;
700             i++;
701         }
702     }
703     returnCode = 0;
704     release(&ptable.lock);
705     return returnCode;
706 }

```

Then we also add new fields of `tickets`, `inuse` and `ticks` to the `proc` structure in file `proc.h`. This will also be helpful when populating the `pstat` structure.

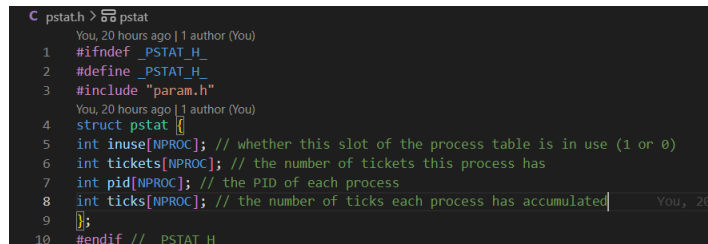


```

C proc.h > DEFAULT_BUDGET
70 // Process original priority
71 // Lottery Scheduling
72 int tickets; // Process tickets
73 int inuse; // Process in use
74 int ticks; // Process ticks
75 };
76

```

The `pstat` structure is defined in the file `pstat.h` file.



```

C pstat.h > pstat
You, 20 hours ago | 1 author (You)
1 #ifndef _PSTAT_H_
2 #define _PSTAT_H_
3 #include "param.h"
4 struct pstat {
5     int inuse[NPROC]; // whether this slot of the process table is in use (1 or 0)
6     int tickets[NPROC]; // the number of tickets this process has
7     int pid[NPROC]; // the PID of each process
8     int ticks[NPROC]; // the number of ticks each process has accumulated
9 };
10 #endif // _PSTAT_H_

```

Algorithm to implement the Lottery Scheduling algorithm we would need a random number generator. We use the C-program for MT19937: Real number version provides us with `sgenrand()`, `genrand()` and `random_at_most()` functions. To do this we add the provided `rand.c` and `rand.h` file. We also add the `rand.o` to the `OBJS` section of `Makefile` to compile the random number generator.

```
C rand.h > random_at_most(long)
You, 20 hours ago | 1 author (You)
1 void sgenrand(unsigned long);
2 long genrand(void);
3 long random_at_most(long);
```

```
M Makefile
30 rand.o
```

As per the question, we the tickets to one and ticks to zero for each process when it is initialize. To do this we edit the `allocproc()` as follows:

```
C proc_lottery_v02.c > allocproc(void)
95 p->state = EMBRYO;
96 p->pid = nextpid++;
97 p->tickets = 1; // By default, each process gets 1 ticket on creation.
98 p->ticks = 0; // By default, each process gets 0 ticks on creation.
99 // settickets(1); // By default, each process gets 1 ticket on creation
```

We also assign, for every child process should have same number of tickets as parent process. To do this we edit the `fork()` function as follows:

```
C proc_lottery_v02.c > fork(void)
210 np->parent = curproc;
211 np->tickets = curproc->tickets; // Copy the number of tickets from the parent
212 *np->tf = *curproc->tf;
```

We also initialize the seed which will be used for random number generator.

```
C proc_lottery_v02.c > fork(void)
210 np->parent = curproc;
211 np->tickets = curproc->tickets; // Copy the number of tickets from the parent
212 *np->tf = *curproc->tf;
```

Algorithm:

```
C proc_lottery_v02.c > scheduler(void)
371 // calculate the total number of tickets
372 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
373     if(p->state != RUNNABLE)
374         continue;
375     ticketSum += p->tickets; // Count the number of tickets of the RUNNABLE processes
376 }
377
378 // Generate a random number between 0 and the total number of tickets to determine which process to run in this lottery round
379 // Change the seed
380 // seed++;
381 sgenrand(seed++);
382 // cprintf("Random Number: %d at time: %d\n", random_at_most(10), uptime());
383
384 uint lotteryNumber = random_at_most(ticketSum);
385
386 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
387     if(p->state != RUNNABLE)
388         continue;
389     // We will use the count variable to keep track of the number of tickets we have counted so far
390     // Then we will compare the count to the random number generated
391     // If the count is less than the random number, we will continue to the next process
392     // If the count is greater than the random number, we will run the current process
393     count += p->tickets;
394     // This is lottery as the process with the most tickets will have the highest chance of being chosen
395     if(count < lotteryNumber){
396         continue;
397     }
398     // Switch to chosen process. It is the process's job
399     // to release ptable.lock and then reacquire it
400     // before jumping back to us.
401     c->xproc = p;
402     switchvm(p);
403     p->state = RUNNING;
404
405     switch(&c->scheduler, p->context);
406     switchvm();
407     // If the process is not completed, update its tickets
408     if(p->state == RUNNABLE && p->tickets < MAXTICKET){
409         p->tickets = p->tickets + 1;
410         // cprintf("INCREMENTED\n");
411     }
412     // Process is done running for now.
413     // It should have changed its p->state before coming back.
414     c->xproc = 0;
415     break;

```

- The algorithm is simple, since each process has an initialized ticket of one, we sum all the tickets of `RUNNABLE` processes. After getting the count, we generate a random number between 0 and the sum. This sum will be used to generate a range for the max random number.
- Then we iterate over all processes and maintain the count of tickets and as soon as the count exceeds over the ticket of recent process, that process gets to run. As a result, the process which has larger number of tickets has the higher probability to run.
- If the process has exited, it has left the system. If the process is still in `RUNNABLE` state, it means it is compute intensive and thus we increase its ticket by one.
- We maintain a `MAXTICKET` count of 15.
- We edit the `pstat` structure each time the scheduler is run so we maintain the latest information.

```
C proc_lottery_v02.c > scheduler(void)
363 // fill the pstat with the processes
364 int i = 0;
365 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
366     pstat.inuse[i] = (p->state == UNUSED) ? 0 : 1;
367     pstat.tickets[i] = p->tickets;
368     pstat.pid[i] = p->pid;
369     pstat.ticks[i] = p->ticks;
370 }
```

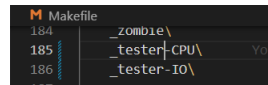
III. Part 03: Comparison and BenchMarking

Final Files: *tester-IO.c* and *tester-CPU.c*

We can compare the performance of the default Round Robin, MLFQ, and Lottery schedulers in xv6 by measuring the performance of each scheduler in terms of throughput and completion time.

We take two tester files using the system call `fork()`. One file will check the CPU and other will check the IO .

To do this we write the code and add the name in `Makefile` under `UPROGS` flag.



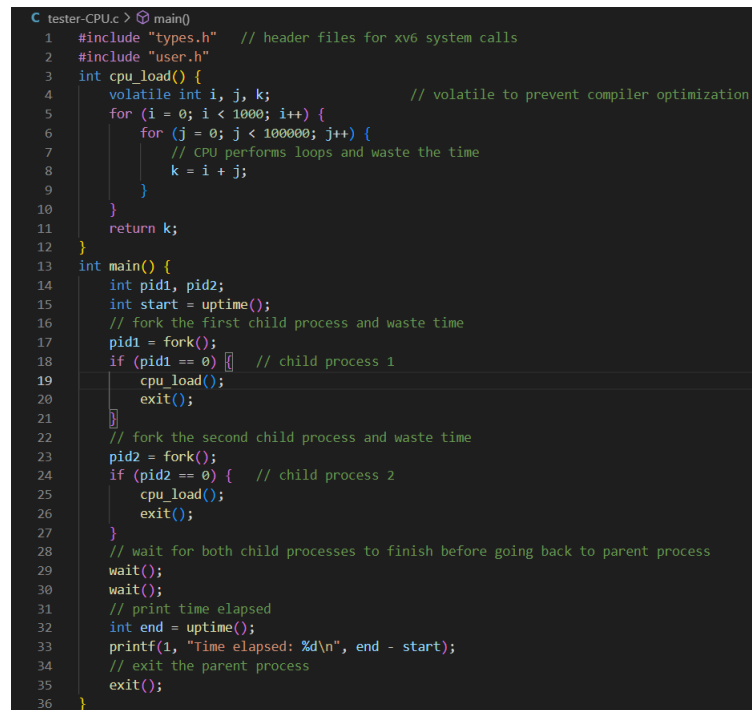
```

184 _zombie\
185 _tester-CPU\
186 _tester-IO\

```

Test Case 1: CPU-Bound Processes

File: `tester-CPU.c`, the definition is:



```

1 #include "types.h" // header files for xv6 system calls
2 #include "user.h"
3 int cpu_load() {
4     volatile int i, j, k; // volatile to prevent compiler optimization
5     for (i = 0; i < 1000; i++) {
6         for (j = 0; j < 100000; j++) {
7             // CPU performs loops and waste the time
8             k = i + j;
9         }
10    }
11    return k;
12 }
13 int main() {
14     int pid1, pid2;
15     int start = uptime();
16     // fork the first child process and waste time
17     pid1 = fork();
18     if (pid1 == 0) // child process 1
19         cpu_load();
20     exit();
21 }
22 // fork the second child process and waste time
23 pid2 = fork();
24 if (pid2 == 0) // child process 2
25     cpu_load();
26 exit();
27 }
28 // wait for both child processes to finish before going back to parent process
29 wait();
30 wait();
31 // print time elapsed
32 int end = uptime();
33 printf(1, "Time elapsed: %d\n", end - start);
34 // exit the parent process
35 exit();
36 }

```

Here, we will create CPU-bound processes that performs large number of calculations (addition) for each process. We create two processes and check the time taken by the program in each of the scheduler.

Test Case 2: I/O-Bound Processes

```

C tester-IO.c > main()
1  #include "types.h" // header files for xv6 system calls
2  #include "user.h"
3  int main() {
4      int i, j;
5      int start = uptime();
6      // create 15 child processes to put the IO load on the cpu
7      for (i = 0; i < 15; i++) {
8          if (fork() == 0) { // child process
9              for (j = 0; j < 100000; j++) {
10                 // just waste time
11             }
12             exit();
13         }
14     }
15     // wait for all child processes to finish
16     for (i = 0; i < 15; i++) {
17         wait();
18     }
19     // Print time elapsed
20     int end = uptime();
21     printf(1, "Time elapsed: %d\n", end - start);
22     // exit the parent process
23     exit();
24 }

```

Here, we will create IO-bound processes that creates a large number of child processes (using `fork()`) for each process. We check the time taken by the program in each of the scheduler. When the child processes are created they write to memory which is an IO operation.

Results:

Overall, we can evaluate the effectiveness of the given schedulers in different scenarios. The time taken in each scenario is:

1. Default (Round Robin) Scheduler:

```

gautam-ahuja@LAPTOP-FV7627LB:~/.../cs1217-lab-2-julius-stabs-back-1$ make qemu-nox
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
Team Julius-Stabs-Back's Schedulers
init: starting sh
$ tester-CPU
Time elapsed: 169
$ tester-IO
Time elapsed: 4
$ |

```

The default scheduler takes 169 ticks for CPU and 4 ticks for IO operation.

2. MLFQ Scheduler:

```

gautam-ahuja@LAPTOP-FV7627LB:~/.../cs1217-lab-2-julius-stabs-back-1$ make qemu-nox
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
Team Julius-Stabs-Back's Schedulers
init: starting sh
$ tester-CPU
Time elapsed: 161
$ tester-IO
Time elapsed: 6

```

The default scheduler takes 161 ticks for CPU and 6 ticks for IO operation.

3. Lottery Scheduler:

```

gautam-ahuja@LAPTOP-FV7627LB:~/.../cs1217-lab-2-julius-stabs-back-1$ make qemu-nox
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
Team Julius-Stabs-Back's Schedulers
init: starting sh
$ tester-CPU
Time elapsed: 632
$ tester-IO
Time elapsed: 3

```

The default scheduler takes 632 ticks for CPU and 3 ticks for IO operation.

From above results, we can see that for CPU intense operation the MLFQ scheduler is the best option and for IO intense operations, the Lottery scheduler is the best option.